

Simplifying iOS Research: Booting iOS on QEMU

@alephsecurity @JonathanAfek

github.com/alephsecurity
alephsecurity.com

Aleph, security research by **HCL AppScan**

How current iOS research is done

- Third party iOS emulator on a remote server
- Development fused iPhone
- Checkm8 demoted iPhone (@axi0mX)
- Off the shelf iPhone - jailbroken
- Off the shelf iPhone - no jailbreak

iPhone panic log

```
1 {"bug_type": "210", "timestamp": "2019-01-31 00:00:31.72 +0100", "os_version": "iPhone OS 11.4.1  
 (15G77)", "incident_id": "C8D49F0C-DBCD-4217-848A-4E3409C968D2"}  
2 {  
3     "build": "iPhone OS 11.4.1 (15G77)",  
4     "product": "iPhone9,4",  
5     "kernel": "Darwin Kernel Version 17.7.0: Mon Jun 11 19:06:26 PDT 2018; root:xnu-4570.70.24~3/RELEASE_ARM64_T8010",  
6     "incident": "C8D49F0C-DBCD-4217-848A-4E3409C968D2",  
7     "crashReporterKey": "bbd821d8854956ba61dd35c63d4c25144a58ad9b",  
8     "date": "2019-01-31 00:00:28.39 +0100",  
9     "panicString": "panic(cpu 0 caller 0xfffffff01b3f8cac): Unaligned kernel data abort. (saved state:  
    0xfffffff024a0b290)\n\t x0: 0xfffffff006405507 x1: 0xfffffff0005fb471 x2: 0x0000000000000001 x3:  
    0x0000000000000000\n\t x4: 0x00000000fffffff x5: 0x0000000000000001 x6: 0x0000000000000000 x7:  
    0xfffffff01ace5a1c\n\t x8: 0xfffffff0005fb470 x9: 0x0000000000000001 x10: 0x0000000000001071ae x11:  
    0x0000000000001071ad\n\t x12: 0x0000000000000000 x13: 0x0000000000000000 x14: 0x0000000000000000 x15:  
    0x0000000000000000\n\t x16: 0xfffffff01b70a014 x17: 0x0000000000000000 x18: 0xfffffff01b2e1000 x19:  
    0xfffffff0064054ff\n\t x20: 0x0000000000000003 x21: 0x0000000000000001 x22: 0xfffffff01b825000 x23:  
    0xfffffff01b2f87e8\n\t x24: 0x0000000000000000a31 x25: 0xfffffff007474000 x26: 0xfffffff0005fb470 x27:  
    0xfffffff0003a5c00\n\t x28: 0x0000000000000000 fp: 0xfffffff024a0b640 lr: 0xfffffff01b2f87a0 sp:  
    0xfffffff024a0b5e0\n\t pc: 0xfffffff01b301c7c cptr: 0xa040304 esr: 0x96000021 far:  
    0xfffffff006405507\n\t Debugger message: panic\n\t Memory ID: 0x1\n\t OS version: 15G77\n\t Kernel version: Darwin Kernel  
    Version 17.7.0: Mon Jun 11 19:06:26 PDT 2018; root:xnu-4570.70.24~3/RELEASE_ARM64_T8010\n\t KernelCache UUID:  
    C50E403D58F345EADBFECFC458171791\n\t niBoot version: iBoot-4076.70.15\n\t insecure boot?: YES\n\t Paniclog version: 9\n\t Kernel  
    slide: 0x0000000014200000\n\t Kernel text base: 0xfffffff01b204000\n\t Epoch Time: sec usec\n\t Boot :  
    0x5c522a5b 0x0008dae2\n\t Sleep : 0x00000000 0x00000000\n\t Wake : 0x00000000 0x00000000\n\t Calendar: 0x5c522c57  
    0x00026e6e\n\t Panicked task 0xfffffff0003e1ce8: 27417 pages, 234 threads: pid 0: kernel_task\n\t Panicked thread:  
    0xfffffff0005fb470, backtrace: 0xfffffff024a0aa90, tid: 401\n\t lr: 0xfffffff01b3f96e8 fp:  
    0xfffffff024a0abd0\n\t lr: 0xfffffff01b2e15f4 fp: 0xfffffff024a0abe0\n\t lr: 0xfffffff01b313a44 fp:  
    0xfffffff024a0af0\n\t lr: 0xfffffff01b313dd4 fp: 0xfffffff024a0afb0\n\t lr: 0xfffffff01b313c00 fp:  
    0xfffffff024a0af0d\n\t lr: 0xfffffff01b3f8cac fp: 0xfffffff024a0b130\n\t lr: 0xfffffff01b3f9dfc fp:  
    0xfffffff024a0b270\n\t lr: 0xfffffff01b2e15f4 fp: 0xfffffff024a0b280\n\t lr: 0xfffffff01b301c7c fp:  
    0xfffffff024a0b640\n\t lr: 0xfffffff01b2f87a0 fp: 0xfffffff024a0b690\n\t lr: 0xfffffff01b2ffa0 fp:  
    0xfffffff024a0b770\n\t lr: 0xfffffff01b306be8 fp: 0xfffffff024a0b850\n\t lr: 0xfffffff01b339988 fp:  
    0xfffffff024a0bb8e0\n\t lr: 0xfffffff01b634b84 fp: 0xfffffff024a0b950\n\t lr: 0xfffffff01b654448 fp:  
    0xfffffff024a0b980\n\t lr: 0xfffffff01b654124 fp: 0xfffffff024a0bab0\n\t lr: 0xfffffff01b654e1c fp:  
    0xfffffff024a0bb30\n\t lr: 0xfffffff01b652ea8 fp: 0xfffffff024a0bc90\n\t lr: 0xfffffff01b2ec500 fp:  
    0x0000000000000000\n\t",  
10    "panicFlags": "0x2",  
11    "otherString": "\n** Stackshot Succeeded ** Bytes Traced 202128 **\n",  
12    "memoryStatus":  
        {"compressorSize": 0, "compressions": 0, "decompressions": 0, "busyBufferCount": 0, "pageSize": 16384, "memoryPressure": false,  
        "memoryPages":  
            {"active": 46523, "throttled": 0, "fileBacked": 69049, "wired": 32283, "purgeable": 1282, "inactive": 14971, "free": 48793,  
            "speculative": 29929}},  
        }  
    }
```

Jonathan Afek

- Aleph Research group manager at HCL/AppScan
- 15 years of experience in security research and low level development including vulnerability research, Linux kernel, storage systems, WiFi systems and FW, security systems and more

iOS on QEMU work done by
@zhuowei (Worth Doing Badly)

QEMU

From Wikipedia, the free encyclopedia

QEMU (short for **Quick EMULATOR**)^[2] is a [free and open-source emulator](#) that performs [hardware virtualization](#).

Past Research - Worth Doing Badly (@zhuowei)

- Chosen version is iPhone X iOS 12 beta 4
 - Extracted the kernel image and the device tree from the software update package
 - Kernel (patched), device tree and the kernel boot arguments were loaded in memory
 - iOS RAMDisk was loaded in memory
 - UART serial support was achieved (output only)
 - Kernel was booted
 - Launchd was executed (no non-Apple executables executed)
-

Past Research - Worth Doing Badly (@zhuowei)

```
BSD root: md0, major 2, minor 0
apfs_vfsop_mountroot:1468: apfs: mountroot called!
apfs_vfsop_mount:1231: unable to root from devvp <ptr> (root_device): 2
apfs_vfsop_mountroot:1472: apfs: mountroot failed, error: 2
hfs: mounted PeaceSeed16A5327f.arm64UpdateRamDisk on device b(2, 0)
: : Darwin Bootstrapper Version 6.0.0: Mon Jul  9 00:39:56 PDT 2018; root:libxpc_execu
boot-args = debug=0x8 kextlog=0xffff cpus=1 rd=md0
Thu Jan  1 00:00:05 1970 localhost com.apple.xpc.launchd[1] <Notice>: Restore environm
```

Goals of our project

- Booting iOS on QEMU ~~with no kernel patches~~
- Supporting hardware (disk, display, touch, sound, multiple CPUs, Interrupt controllers, etc...)
- Supporting different iOS versions
- Conducting iOS security research
- Learning about iOS and QEMU internals

Status of our project

- Booting ~~Secure Monitor~~ and the kernel (~~unpatched~~) on QEMU
 - Executing user-mode apps over launchd
 - Running an interactive bash shell on an iOS kernel on QEMU
 - Running gdb scripts listing and switching (tasks, threads, zones)
 - Allowing access to tfp0 from user mode
 - Mounting full disk images with our own block device driver
 - Running most of the iOS user mode services
 - Connecting to an SSH server running on the iOS system
 - Displaying a textual iOS framebuffer
 - Supporting only on iOS 12.1 for iPhone 6s plus for now
-

```
Jonathans-MBP:demo jonathanafek$ /Users/jonathanafek/src/aleph/xnu-qemu-arm64/aarch64-softmmu/qemu-system-aarch64 -M iPhone6splus-n66-s8000,kernel-filename=fw/kernelcache.release.n66.out,dtb-filename=fw/Firmware/all_flash/DeviceTree.n66ap.im4p.out,tc-filename=static_tc,driver-filename=/Users/jonathanafek/src/aleph/xnu-qemu-arm64-tools/aleph_bdev_drv/bin/aleph_bdev_drv.bin,qc-file-1-filename=/Users/jonathanafek/demo/hfs.big_sec,qc-file-0-filename=/Users/jonathanafek/demo/hfs.big_main,kern-cmd-args="debug=0x8 kextlog=0xffff cpus=1 rd=disk0 serial=2" -cpu max -m 6G -serial mon:stdio -nographic -S -s
```

[1 bash] [2 bash] [3 bash] [4 bash] [5 bash] [6 bash]

Wednesday 13:06 29/01/2020

```
(reverse-i-search)`ramf': /Users/jonathanafek/src/aleph/xnu-qemu-arm64/aarch64-softmmu/qemu-system-aarch64 -M iPhone6splus-n66-s8000,kernel-filename=fw/kernelcache.release.n66.out,dtb-filename=fw/Firmware/all_flash/DeviceTree.n66ap.im4p.out,tc-filename=static_tc,driver-filename=/Users/jonathanafek/src/aleph/xnu-qemu-arm64-tools/aleph_bdev_drv/bin/aleph_bdev_drv.bin,qc-file-1-filename=/Users/jonathanafek/demo/hfs.big_sec,qc-file-0-filename=/Users/jonathanafek/demo/hfs.big_main,xnu-ramfb=on,kern-cmd-args="debug=0x8 kextlog=0xffff cpus=1 rd=disk0 serial=2" -cpu max -m 6G -serial mon:stdio
```

[1 bash] [2 bash] [3 bash] [4 bash] [5 bash] [6 bash]

Wednesday 13:08 29/01/2020

Agenda

- Past public research on iOS on QEMU
 - **iOS kernel boot process**
 - Execution of non-Apple executables with Trust Cache
 - Bash execution with launchd
 - tfp0
 - SSH connection
 - Textual framebuffer
 - Block device driver
 - No secure monitor
 - Launchd craziness
 - Next steps
-

iOS kernel boot process

- Start booting the kernelcache code in EL1 as done by @zhuowei
- Crash on SMC instruction (Secure Monitor Call)

Register group: general

x0 0x800	2048	x1 0x470a5000	1191858176	x2 0x0	0
x6 0xff	255	x7 0xad0	2768	x8 0x40000000	1073741824
x12 0x14	1	x13 0x1000000	16777216	x14 0x3000000	50331648
x18 0x0	0	x19 0xfffffffff0075d1000	-68595937280	x20 0xfffffffff0070a5088	-68601360248
x24 0xf4240	1000000	x25 0xfffffffff00707d0000	-68601524224	x26 0xfffffffff007078440	-68601543616
x30 0xfffffff0071b74a0	-68600236896	x27 0xfffffffff00767fe10	0xfffffff00767fe1pc	x28 0xfffffff0070a7d3c	0xfffffff0070a7d7c
IMVFR6_EL1_RESERVED 0x0	0	MVFR4_EL1_RESERVED 0x0	0	ID_AA64PFR1_EL1 0x0	0
ID_AA64PFR5_EL1_RESERVED 0x0	0	ID_AA64PFR6_EL1_RESERVED 0x0	0	ID_AA64PFR7_EL1_RESERVED 0x0	0
ID_AA64DFR3_EL1_RESERVED 0x0	0	ID_AA64AFR1_EL1 0x0	0	ID_AA64AFR2_EL1_RESERVED 0x0	0
ID_AA64ISAR5_EL1_RESERVED 0x0	0	ID_AA64ISAR3_EL1_RESERVED 0x0	0	ID_AA64ISAR1_EL1 0x0	0
ID_AA64AFR0_EL1 0x0	0	ID_AA64MMFR1_EL1 0x0	0	ID_AA64MMFR2_EL1_RESERVED 0x0	0
ID_AA64MMFR6_EL1_RESERVED 0x0	0	ID_AA64MMFR0_EL1 0x1124	4388	ID_AA64MMFR5_EL1_RESERVED 0x0	0
REVIDR_EL1 0x0	0	SCTLR 0x3454593d	877943101	ACTLR_EL1 0x0	0
SCTRLR_EL2 0x0	0	HSTR_EL2 0x0	0	CPTLR_EL2 0x0	0
IMDCR_EL2 0x0	0	MDCR_EL3 0x0	0	SCTRLR_EL3 0x30d5180d	819271693
PMCR_EL0 0x41000000	1090519040	PMCNTENCLR_EL0 0x0	0	PMCNTENSET_EL0 0x0	0
PMCCNTR_EL0 0x0	0	PMCEIDI_EL0 0x0	0	TIBR0_EL1 0x10000a7a40000	281477789253632
MAIR_EL1 0x44f00bb44ff	4737361200383	TIBR0_EL2 0x0	0	TCR_EL2 0x0	0
TIBR0_EL3 0x410000c000	279172923392	MAIR_EL2 0x0	0	TCR_EL3 0x1a511	107793
L2CTLR_EL1 0x0	0	L2ECTLR_EL1 0x0	0	DACR32_EL2 0x0	0
SP_EL0 0x0	0	VBAR_EL2 0x0	0	SP_EL1 0xfffffff007688000	-68595187712
ISPSCR_TRO 0x0	0	SPSR_ABТ 0x0	0	FPCR 0x0	0

B-> [0xfffffff0070a7d3c] smc #0x11

[0xfffffff0070a7d40] ret

0xfffffff0070a7d44 .inst 0x00000000 ; undefined

0xfffffff0070a7d48 .inst 0x00000000 ; undefined

0xfffffff0070a7d4c .inst 0x00000000 ; undefined

0xfffffff0070a7d50 .inst 0x00000000 ; undefined

0xfffffff0070a7d54 .inst 0x00000000 ; undefined

0xfffffff0070a7d58 .inst 0x00000000 ; undefined

0xfffffff0070a7d5c .inst 0x00000000 ; undefined

0xfffffff0070a7d60 .inst 0x03020100 ; undefined

0xfffffff0070a7d64 .inst 0x07060504 ; undefined

0xfffffff0070a7d68 add w8, w8, w10, lsl #2

0xfffffff0070a7d6c .inst 0x0f0e000c ; undefined

0xfffffff0070a7d70 .inst 0x00000000 ; undefined

0xfffffff0070a7d74 .inst 0x00000000 ; undefined

0xfffffff0070a7d78 .inst 0x00000000 ; undefined

0xfffffff0070a7d7c .inst 0x00000000 ; undefined

0xfffffff0070a7d80 tst x1, x1

0xfffffff0070a7d84 b.mi 0xfffffff0070a7e20 // b.first

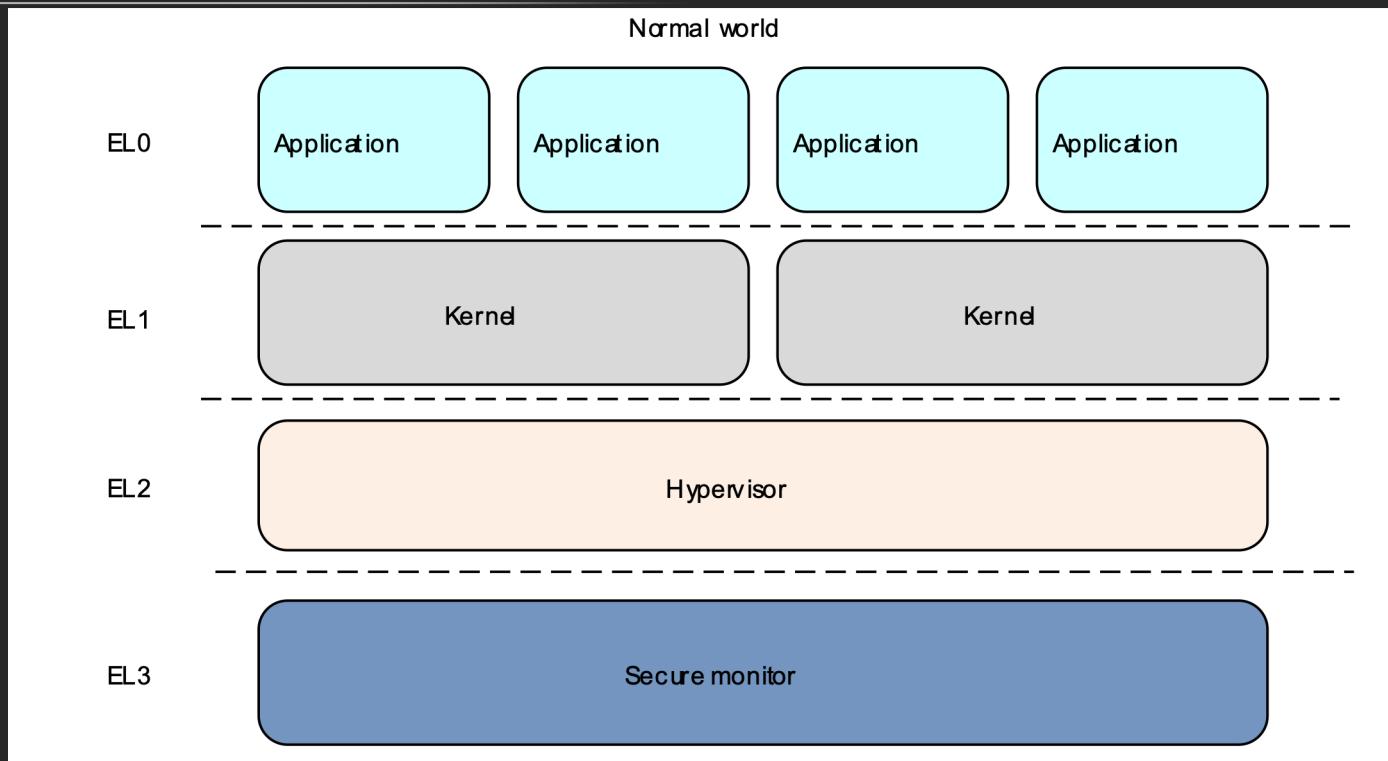
0xfffffff0070a7d88 b.eq 0xfffffff0070a7dec // b.none

0xfffffff0070a7d8c stp x29, x30, [sp, #-16]!

0xfffffff0070a7d90 mov x29, sp

remote Thread 1 In:

iOS kernel boot process



from <https://developer.arm.com/docs/den0024/a/fundamentals-of-armv8>

iOS kernel boot process

- Secure Monitor starts execution at boot in EL3
- It resides in a secure memory location inaccessible from EL1 (kernel code)
- It services SMC calls from the kernel (similar to how system calls from user apps to the kernel are serviced)
- It is responsible for KPP (Kernel Patch Protection) in our system

iOS kernel boot process

- Load the Secure Monitor image for iOS 12.1 for iPhone 6s plus in the secure memory
 - Load the secure monitor boot arguments in the secure memory with the relevant values
 - Start executing the secure monitor at its entry point in EL3
 - The secure monitor finishes its boot sequence and jumps to the kernel at its entry point in EL1
 - The kernel finishes its boot seq and uses the secure monitor to service its SMC calls
-

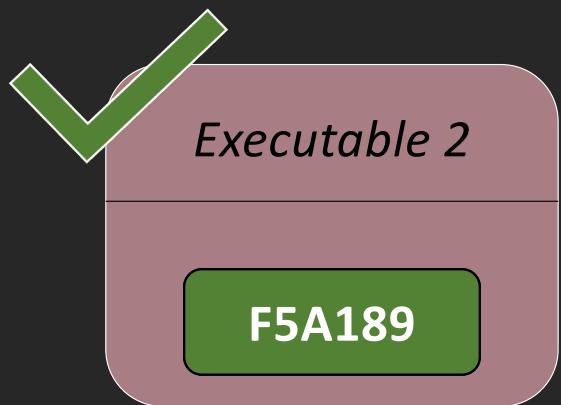
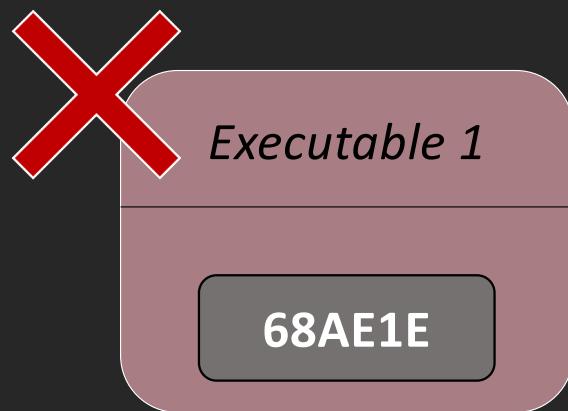
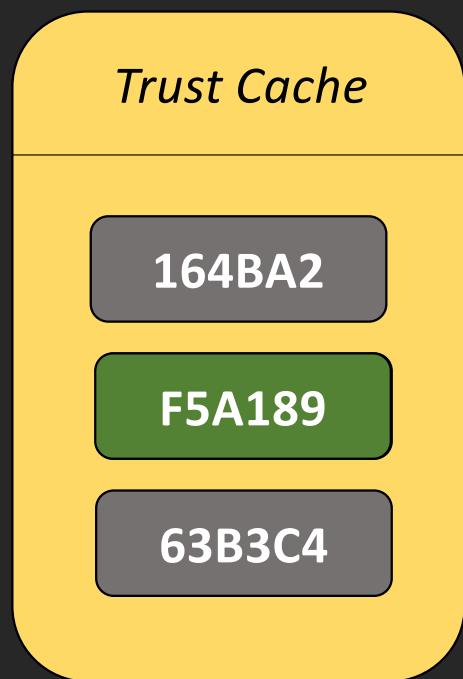
iOS kernel boot process

And it works!

Agenda

- Past public research on iOS on QEMU
 - iOS kernel boot process
 - **Execution of non-Apple executables with Trust Cache**
 - Bash execution with launchd
 - tfp0
 - SSH connection
 - Textual framebuffer
 - Block device driver
 - No secure monitor
 - Launchd craziness
 - Next steps
-

Trust Cache



Trust Cache

- iOS has 3 different types of trust caches
 - A list of hardcoded hashes approved in the kernelcache
 - A dynamic trust cache that can be loaded at runtime from a file
 - A static trust cache in memory pointed from the device tree

Trust Cache

Device tree

From Wikipedia, the free encyclopedia

In computing, a **device tree** (also written **devicetree**) is a **data structure** describing the hardware components of a particular computer so that the operating system's **kernel** can use and manage those components, including the **CPU** or CPUs, the **memory**, the **buses** and the **peripherals**.

Trust Cache

- By reversing this code and other relevant code we were able to understand the input structure and construct our own static TC

The screenshot shows a debugger interface with two panes. The left pane displays assembly code with various registers and memory locations highlighted in blue and orange. The right pane displays corresponding C code with variable names and function definitions.

Assembly Code (Left Pane):

```
...ff0071b5c9c e0 07 9f 1a    cset    w0,ne
...ff0071b5ca0 04 00 00 14    b       LAB_ffffffff0071b5cb0

LAB_ffffffff0071b5ca4
...ff0071b5ca4 9f fe 9f c8    stlr    xzr,[x20]>=DAT_ffffffff00764bb80
...ff0071b5ca8 85 1c 00 94    bl      FUN_ffffffff0071bcebc
...ff0071b5cac e0 03 00 32    mov     w0,#0x1

LAB_ffffffff0071b5cb0
...ff0071b5cb0 fd 7b 42 a9    ldp    x29>local_10,x30,[sp], #0x20
...ff0071b5cb4 f4 4f 41 a9    ldp    x20,x19,[sp], #local_20
...ff0071b5cb8 f6 57 c3 a8    ldp    x22,x21,[sp], #0x30
...ff0071b5cbc c0 03 5f d6    ret

***** FUNCTION *****
ulonglong __cdecl _pmap_lookup_in_static_trust_cache(by...
    ulonglong x0:8    <RETURN>
    byte *    x0:8    pbParm1_input_hash
    ulonglong x8:8    uVar5_hash_table_part_size
    byte *    x9:8    pbVar5_pointer_to_current_part_of_hash...
    ulonglong x11:8   uVar5_index_hash_table
    undefined8 Stack[-0x10]:8 local_10
    undefined8 Stack[-0x20]:8 local_20
    undefined8 Stack[-0x30]:8 local_30
    _pmap_lookup_in_static_trust_cache
    ...ff0071b5cc0 ff c3 00 d1    sub    sp,sp,#0x30
    ...ff0071b5cc4 f4 4f 01 a9    stp    x20,x19,[sp], #local_20
```

C Code (Right Pane):

```
2  ulonglong _pmap_lookup_in_static_trust_cache(byte *pbParm1_input_hash)
3
4  long long lVar1;
5  byte *pbVar2;
6  byte bVar3;
7  ushort uVar4;
8
9  ulonglong uVar5_hash_table_part_size;
10 byte *pbVar5_pointer_to_current_part_of_hash_table;
11 ulonglong uVar5;
12 ulonglong uVar5_index_hash_table;
13 ulonglong uVar6;
14 uint uVar7;
15 uint uVar8;
16
17 if (DAT_ffffffff00707de28 != 0) {
18     uVar5_hash_table_part_size = (ulonglong)*(uint *) (DAT_ffffffff00707de28 + 0x14);
19     if ((*uint *) (DAT_ffffffff00707de28 + 0x14) != 0) {
20         pbVar5_pointer_to_current_part_of_hash_table = (byte *) (DAT_ffffffff00707de28 + 0x18);
21         do {
22             uVar5_index_hash_table = uVar5_hash_table_part_size >> 1;
23             uVar7 = (uint)pbParm1_input_hash;
24             bVar3 = pbVar5_pointer_to_current_part_of_hash_table[uVar5_index_hash_table * 0x16];
25             uVar8 = (uint)bVar3;
26             if (pbParm1_input_hash == bVar3) {
27                 uVar7 = (uint)pbParm1_input_hash[1];
28                 uVar8 = (uint)pbVar5_pointer_to_current_part_of_hash_table
29                             [uVar5_index_hash_table * 0x16 + 1];
30             if (pbParm1_input_hash[1] ==
31                 pbVar5_pointer_to_current_part_of_hash_table[uVar5_index_hash_table * 0x16 + 1]) {
32                 uVar7 = (uint)pbParm1_input_hash[2];
33                 uVar8 = (uint)pbVar5_pointer_to_current_part_of_hash_table
34                             [uVar5_index_hash_table * 0x16 + 2];
35             if (pbParm1_input_hash[2] ==
36                 pbVar5_pointer_to_current_part_of_hash_table[uVar5_index_hash_table * 0x16 + 2]) {
37                 uVar7 = (uint)pbParm1_input_hash[2];
38             }
39         }
40     }
41 }
```

Trust Cache

- Which Apple released the source code for

```
DTEntry memory_map;
MemoryMapFileInfo *trustCacheRange;
unsigned int trustCacheRangeSize;
int err;

err = DTLookupEntry(NULL, "chosen/memory-map", &memory_map);
assert(err == kSuccess);

err = DTGetProperty(memory_map, "TrustCache", (void**)&trustCacheRange
if (err == kSuccess) {
    assert(trustCacheRangeSize == sizeof(MemoryMapFileInfo));

    segEXTRADATA = phystokv(trustCacheRange->paddr);
    segSizeEXTRADATA = trustCacheRange->length;

    arm_vm_page_granular_RNX(segEXTRADATA, segSizeEXTRADATA, FALSE
}
```

Trust Cache

And it works!



Agenda

- Past public research on iOS on QEMU
 - iOS kernel boot process
 - Execution of non-Apple executables with Trust Cache
 - **Bash execution with launchd**
 - tfp0
 - SSH connection
 - Textual framebuffer
 - Block device driver
 - No secure monitor
 - Launchd craziness
 - Next steps
-

Bash on Launchd

- Mount the RAMDisk image on OSX
- Remove all files in /System/Library/LaunchDaemons/
- Add a single file there for running bash
(com.apple.bash.plist)
- Add the bash executable to the RAMDisk
- Add the bash executable hash to the Trust Cache
- Unmount the RAMDisk and run QEMU

Bash on Launchd

```
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE plist PUBLIC "-//Apple//DTD PLIST 1.0//EN" "http://www.apple.com/DTDs/PropertyList-1.0.dtd">
<plist version="1.0">
<dict>
    <key>EnablePressuredExit</key>
    <false/>
    <key>Label</key>
    <string>com.apple.bash</string>
    <key>POSIXSpawnType</key>
    <string>Interactive</string>
    <key>ProgramArguments</key>
    <array>
        <string>/iosbinpack64/bin/bash</string>
    </array>
    <key>RunAtLoad</key>
    <true/>
    <key>StandardErrorPath</key>
    <string>/dev/console</string>
    <key>StandardInPath</key>
    <string>/dev/console</string>
    <key>StandardOutPath</key>
    <string>/dev/console</string>
    <key>Umask</key>
    <integer>0</integer>
    <key>UserName</key>
    <string>root</string>
</dict>
</plist>
```

Bash on Launchd

- System tries to execute bash
- Logs show missing libraries required for bash

Bash on Launchd

Any ideas?

Bash on Launchd

- The RAMDisk image comes without the dynamic loader cache on it, which is a file that holds most of the common runtime libs for iOS
- Copy this file into the RAMDisk at the correct path from the full disk images

Bash on Launchd

Any ideas?

Bash on Launchd

- Debug `/usr/lib/dyld` (the dynamic loader) which is responsible for loading the dynamic loader cache

Bash on Launchd

```
// map in shared cache to shared region
int fd = openSharedCacheFile();
if ( fd != -1 ) {
    uint8_t firstPages[8192];
    if ( ::read(fd, firstPages, 8192) == 8192 ) {
        dyld_cache_header* header = (dyld_cache_header*)firstP
#if __x86_64__
        const char* magic = (sHaswell ? ARCH_CACHE_MAGIC_H : A
#else
        const char* magic = ARCH_CACHE_MAGIC;
#endif
        if ( strcmp(header->magic, magic) == 0 ) {
            const dyld_cache_mapping_info* const fileMappi
            const dyld_cache_mapping_info* const fileMappi
            shared_file_mapping_np  mappings[header->mappi
            unsigned int mappingCount = header->mappingCount;
```

Bash on Launchd

- Stepping through the execution path with gdb showed the error was in here

```
if (_shared_region_map_and_slide_np(fd, mappingCount, mappings, codeSize)) {
    // successfully mapped cache into shared region
    sSharedCache = (dyld_cache_header*)mappings[0].sfm_address;
    sSharedCacheSlide = cacheSlide;
    dyld::gProcessInfo->sharedCacheSlide = cacheSlide;
    //dyld::log("sSharedCache=%p sSharedCacheSlide=0x%08lX\n", sSharedCache,
    // if cache has a uuid, copy it
    if ( header->mappingOffset >= 0x68 ) {
        memcpy(dyld::gProcessInfo->sharedCacheUUID, header->uuid,
               16);
    }
} else {
    throw "dyld shared cache could not be mapped";
    if ( gLinkContext.verboseMapping )
        dyld::log("dyld: shared cached file could not be mapped");
}
```

Bash on Launchd

- Since we have a kernel debugger in gdb we can step into the system call in the kernel

```
int
shared_region_map_and_slide_np(
    struct proc                                *p,
    struct shared_region_map_and_slide_np_args   *uap,
    __unused int                                     *retvalp)
{
    struct shared_file_mapping_np    *mappings;
    unsigned int                      mappings_count = uap->count;
    kern_return_t                     kr = KERN_SUCCESS;
    uint32_t                          slide = uap->slide;

#define SFM_MAX_STACK    8
    struct shared_file_mapping_np    stack_mappings[SFM_MAX_STACK]

    /* Is the process chrooted?? */
    if (p->p_fd->fd_rdir != NULL) {
        kr = EINVAL;
        goto done;
    }
```

Bash on Launchd

- Stepping through this function we see that the call to `_shared_region_map_and_slide()` is the part that fails

```
kr = _shared_region_map_and_slide(p, uap->fd, mappings_count, mapping  
slide,  
uap->slide_start, uap->slide_size);  
if (kr != KERN_SUCCESS) {  
    return kr;  
}  
  
return kr;
```

Bash on Launchd

- Stepping in that function reveals the error here

```
/* make sure vnode is owned by "root" */
VATTR_INIT(&va);
VATTR_WANTED(&va, va_uid);
error = vnode_getattr(vp, &va, vfs_context_current());
if (error) {
    SHARED_REGION_TRACE_ERROR(
        ("shared_region: %p [%d(%s)] map(%p:'%s'): "
         "vnode_getattr(%p) failed (error=%d)\n",
         (void *)VM_KERNEL_ADDRPERM(current_thread()),
         p->p_pid, p->p_comm,
         (void *)VM_KERNEL_ADDRPERM(vp), vp->v_name,
         (void *)VM_KERNEL_ADDRPERM(vp), error));
    goto done;
}
if (va.va_uid != 0) {
    SHARED_REGION_TRACE_ERROR(
        ("shared_region: %p [%d(%s)] map(%p:'%s'): "
         "owned by uid=%d instead of 0\n",
         (void *)VM_KERNEL_ADDRPERM(current_thread()),
         p->p_pid, p->p_comm,
         (void *)VM_KERNEL_ADDRPERM(vp),
         vp->v_name, va.va_uid));
    error = EPERM;
    goto done;
}
```

Bash on Launchd

- The code validates that the cache file is owned by root
- Mount the RAMDisk image in a different way to allow permission editing
- Copy the cache file and chown to root

Bash on Launchd

And it works!



Agenda

- Past public research on iOS on QEMU
 - iOS kernel boot process
 - Execution of non-Apple executables with Trust Cache
 - Bash execution with launchd
 - **tfp0**
 - SSH connection
 - Textual framebuffer
 - Block device driver
 - No secure monitor
 - Launchd craziness
 - Next steps
-

tfp0

- `task_for_pid()` is a mach trap (like a system call) that allows getting access to another task port on the same host from usermode
- For pid 0 this gives access to the kernel task port
- Access to this port gives full control over the other task (memory read/write, code execution, etc..)

tfp0

```
kern_return_t task_for_pid(struct task_for_pid_args *args)
{
    // ...

    /* Always check if pid == 0 */
    if (pid == 0) {
        (void) copyout((char *)&t1, task_addr, sizeof(mach_port_name_t));
        AUDIT_MACH_SYSCALL_EXIT(KERN_FAILURE);
        return(KERN_FAILURE);
    }

    // ...
}
```

from <https://github.com/Siguza/hsp4/blob/master/README.md>



tfp0

```
task_t convert_port_to_task_with_exec_token(ipc_port_t port, uint32_t *exec_token)
{
    // ...

    if (task == kernel_task && current_task() != kernel_task) {
        ip_unlock(port);
        return TASK_NULL;
    }

    // ...
}
```

from <https://github.com/Siguza/hsp4/blob/master/README.md>



tfp0 (@s1guza)

```
vm_map_remap(
    kernel_map,
    &remap_addr,
    sizeof(task_t),
    0,
    VM_FLAGS_ANYWHERE | VM_FLAGS_RETURN_DATA_ADDR,
    zone_map,
    kernel_task,
    false,
    &dummy,
    &dummy,
    VM_INHERIT_NONE
);
// mach_vm_wire_external as of High Sierra
mach_vm_wire(&realhost, kernel_map, remap_addr, sizeof(task_t), VM_PROT_READ | VM_PROT_WRITE);
ipc_port_t port = ipc_port_alloc_special(ipc_space_kernel);
ipc_kobject_set(port, remap_addr, IKOT_TASK);
realhost.special[4] = ipc_port_make_send(port);
```

from <https://github.com/Siguza/hsp4/blob/master/README.md>

tfp0

- Allocate static virtual memory for the remapped kernel task
- QEMU “MMIO” to redirect reads/writes from/to our fake kernel task to the real kernel task

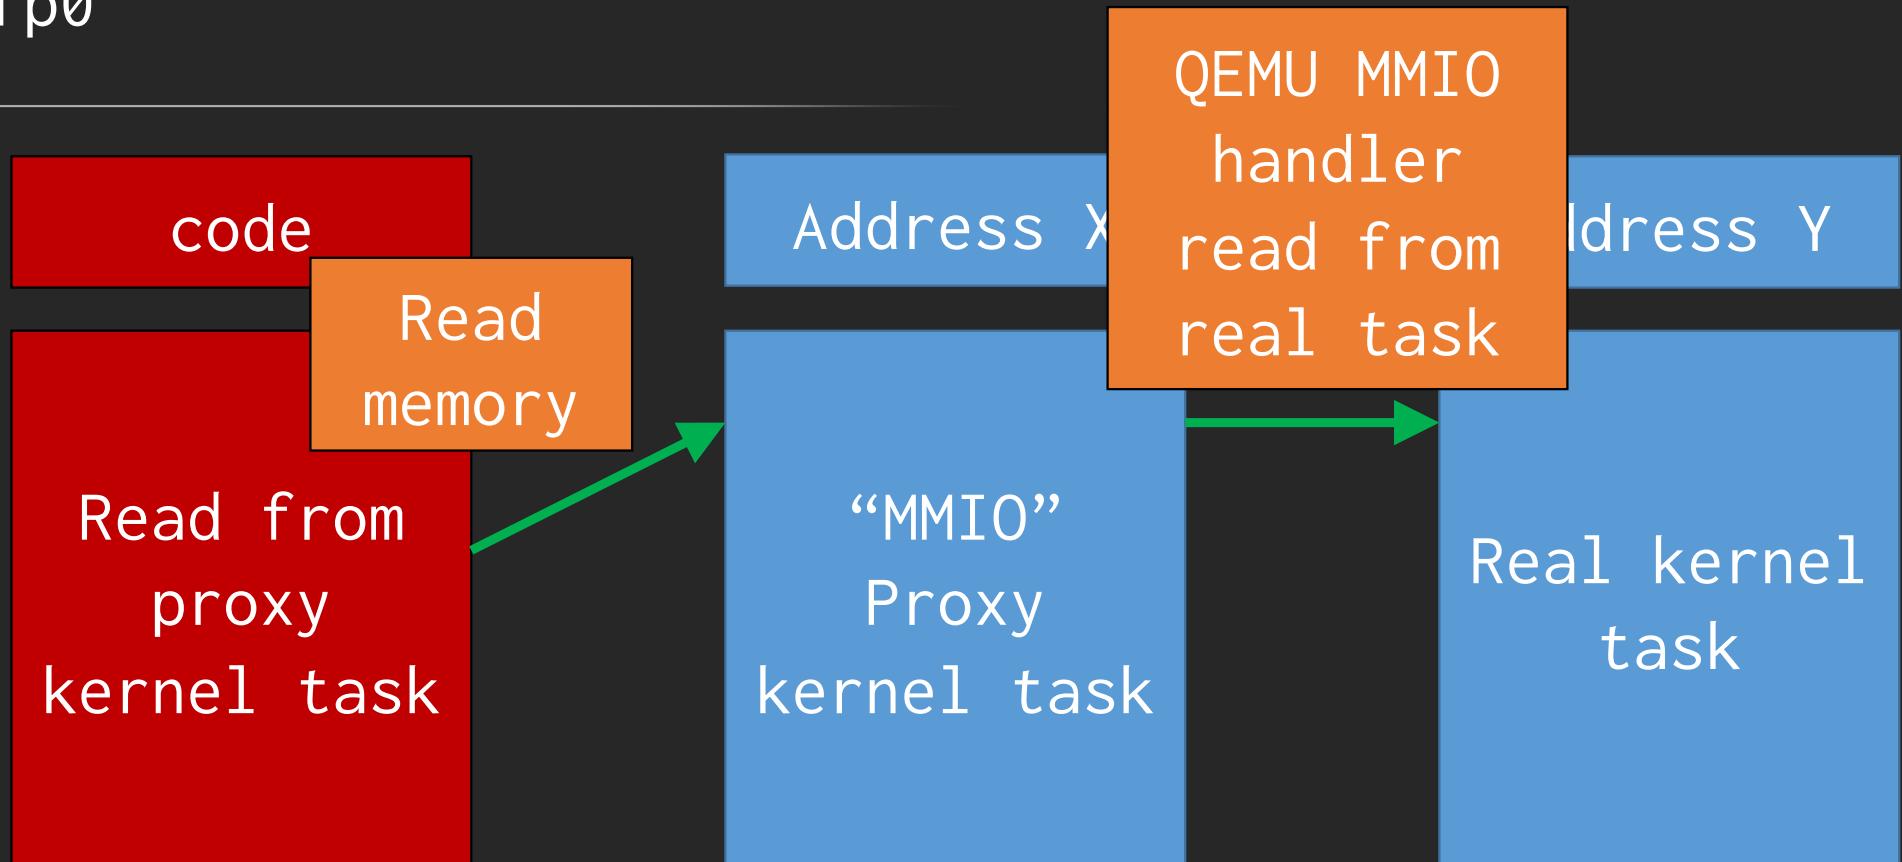
```
ktpp->remap.orig_pa = vtop_mmu(global_kernel_ptr, ktpp->cs);$  
ktpp->remap.pa = remap_task_paddr;$  
ktpp->remap.orig_as = ktpp->as;$  
ktpp->remap.size = KERNEL_TASK_ALLOC_SIZE;$  
  
xnu_dev_remap_create(sysmem, &ktpp->remap, "kernel_task_remap_dev");
```


tfp0

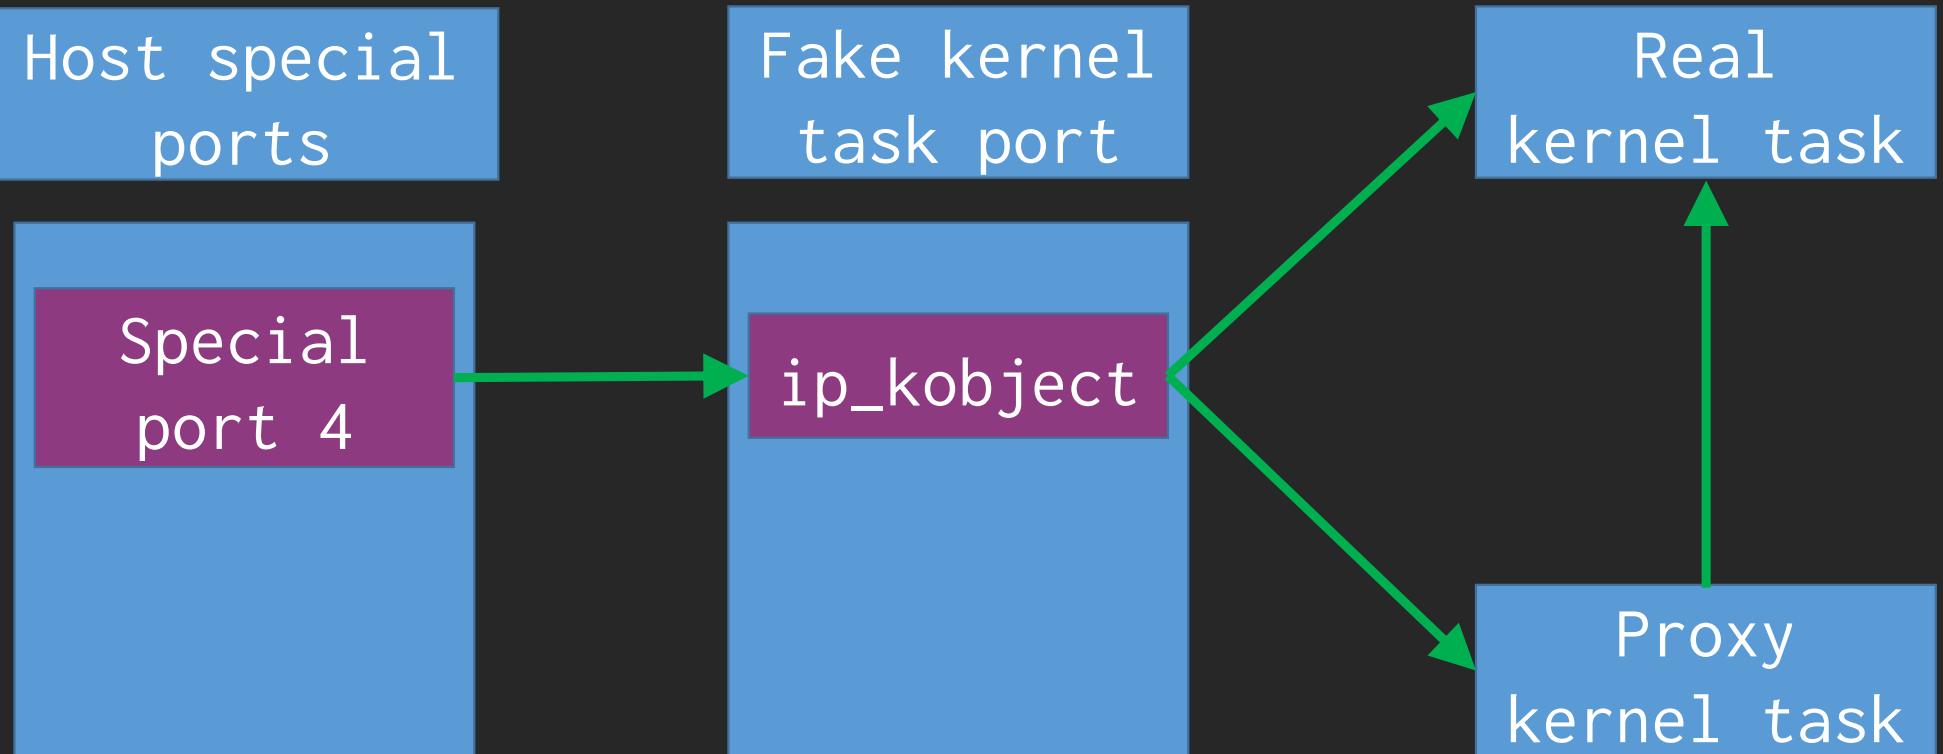
- Copy the port to host special port 4

```
address_space_rw(ktpp->as, vtop_mmu(special_port_4_vaddr, ktpp->cs),$  
    MEMTXATTRS_UNSPECIFIED, (uint8_t *)&fake_port_vaddr,$  
    sizeof(hwaddr), 1);$
```

tfp0



tfp0



tfp0

And it works!

```
-bash-4.4# kmem 0xffffffff005c5c000 256
[*] Reading 256 bytes from 0xffffffff005c5c000
CF FA ED FE 0C 00 00 01 00 00 00 00 0B 00 00 00 |.....|
03 00 00 00 78 00 00 00 02 00 00 00 00 00 00 00 |....x.....|
19 00 00 00 48 00 00 00 5F 5F 4C 49 4E 4B 45 44 |....H...__LINKED|
49 54 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |IT.....|
68 0A 00 00 00 00 00 00 00 10 00 00 00 00 00 00 |h.....|
68 0A 00 00 00 00 00 00 01 00 00 01 00 00 00 00 |h.....|
00 00 00 00 04 00 00 00 02 00 00 00 18 00 00 00 |.....|
00 10 00 00 44 00 00 00 40 14 00 00 28 06 00 00 |....D...@...C...|
CF FA ED FE 0C 00 00 01 00 00 00 00 0B 00 00 00 |.....|
03 00 00 00 78 00 00 00 02 00 00 00 00 00 00 00 |....x.....|
19 00 00 00 48 00 00 00 5F 5F 4C 49 4E 4B 45 44 |....H...__LINKED|
49 54 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |IT.....|
E4 96 00 00 00 00 00 00 00 10 00 00 00 00 00 00 |.....|
E4 96 00 00 00 00 00 00 01 00 00 00 01 00 00 00 |.....|
00 00 00 00 04 00 00 00 02 00 00 00 18 00 00 00 |.....|
00 10 00 00 99 03 00 00 90 49 00 00 54 5D 00 00 |.....I..T]..|
```

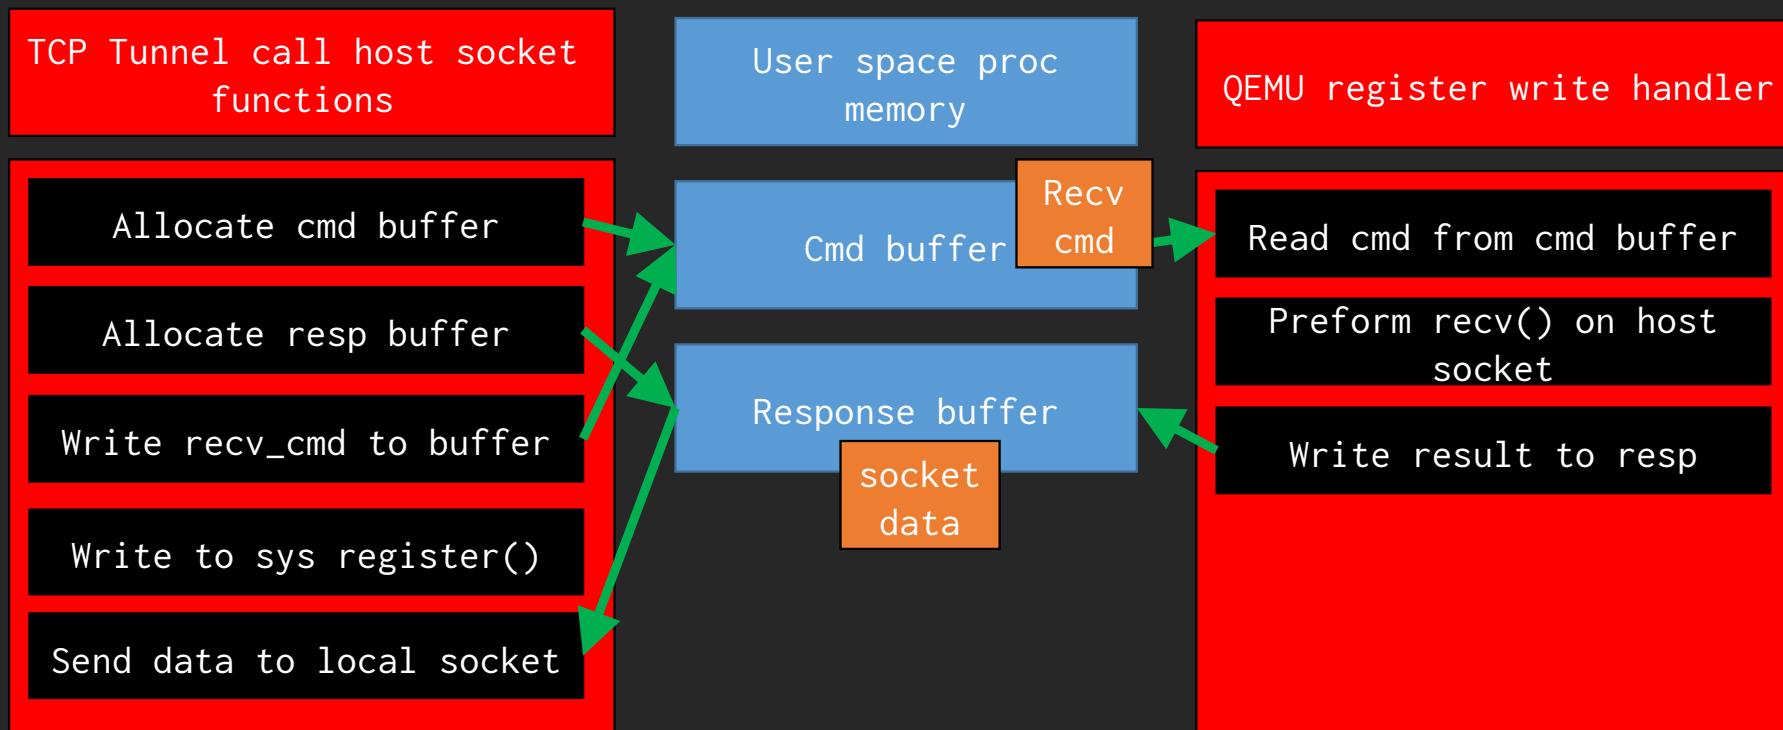
Agenda

- Past public research on iOS on QEMU
 - iOS kernel boot process
 - Execution of non-Apple executables with Trust Cache
 - Bash execution with launchd
 - tfp0
 - **SSH connection**
 - Textual framebuffer
 - Block device driver
 - No secure monitor
 - Launchd craziness
 - Next steps
-

TCP connections and SSH (@levaronsky)

- The “normal” way emulating the networking hardware
- The “normal” way with a driver and an emulated network controller for that driver
- Our way with communicating with the QEMU host from a usermode TCP tunnel proc using QEMU’s register write callbacks

TCP connections and SSH



TCP connections and SSH

- recv() from host socket and send() to guest socket and vice versa
- Very inefficient with a busy loop running endlessly without blocking at all when no traffic is available
- Works smoothly enough for now

```
Jonathans-MBP:~ jonathanafek$ ssh -p 2222 root@127.0.0.1
root@127.0.0.1's password:
[104] Dec 31 16:07:57 lastlog_perform_login: Couldn't stat /var/log/lastlog: No such file or directory
[104] Dec 31 16:07:57 lastlog_openseek: /var/log/lastlog is not a file or directory!
-bash-4.4# export PATH=$PATH:/iosbinpack64/usr/bin:/iosbinpack64/bin:/iosbinpack64/usr/sbin:/iosbinpack64/sbin
-bash-4.4# uname -a
Darwin localhost 18.2.0 Darwin Kernel Version 18.2.0: Tue Oct 16 21:02:23 PDT 2018; root:xnu-4903.222.5~1/RELEASE_ARM64_S8000 iPhone8,2
```

Agenda

- Past public research on iOS on QEMU
 - iOS kernel boot process
 - Execution of non-Apple executables with Trust Cache
 - Bash execution with launchd
 - tfp0
 - SSH connection
 - **Textual framebuffer**
 - Block device driver
 - No secure monitor
 - Launchd craziness
 - Next steps
-

Textual framebuffer (@V3rochka)

- Setup a new ramfb device with our own display parameters on the QEMU host side

```
void xnu_define_ramfb_device(AddressSpace* as, hwaddr ramfb_pa){$  
$  
    DeviceState *fb_dev;$  
    fb_dev = qdev_create(NULL, TYPE_XNU_RAMFB_DEVICE);$  
    qdev_prop_set_uint64(fb_dev, "as", (hwaddr)as);$  
    qdev_prop_set_uint64(fb_dev, "fb_pa", ramfb_pa);$  
    qdev_prop_set_uint32(fb_dev, "fb_size", RAMFB_SIZE);$  
    qdev_prop_set_uint32(fb_dev, "display_cfg.height", V_HEIGHT);$  
    qdev_prop_set_uint32(fb_dev, "display_cfg.width", V_WIDTH);$  
    qdev_prop_set_uint32(fb_dev, "display_cfg.linesize", V_LINESIZE);$  
    qdev_init_nofail(fb_dev);$  
}$
```

Textual framebuffer

- Setup the iOS boot args to configure the textual framebuffer

```
void xnu_get_video_bootargs(void *opaque, hwaddr ramfb_pa){$  
$    video_boot_args* v_bootargs = (video_boot_args*)opaque ;$  
    v_bootargs->v_baseAddr = ramfb_pa;$  
    v_bootargs->v_depth = V_DEPTH;$  
    v_bootargs->v_display = V_DISPLAY;$  
    v_bootargs->v_height = V_HEIGHT;$  
    v_bootargs->v_width = V_WIDTH;$  
    v_bootargs->v_rowBytes = V_LINESIZE;$  
$}  
$
```

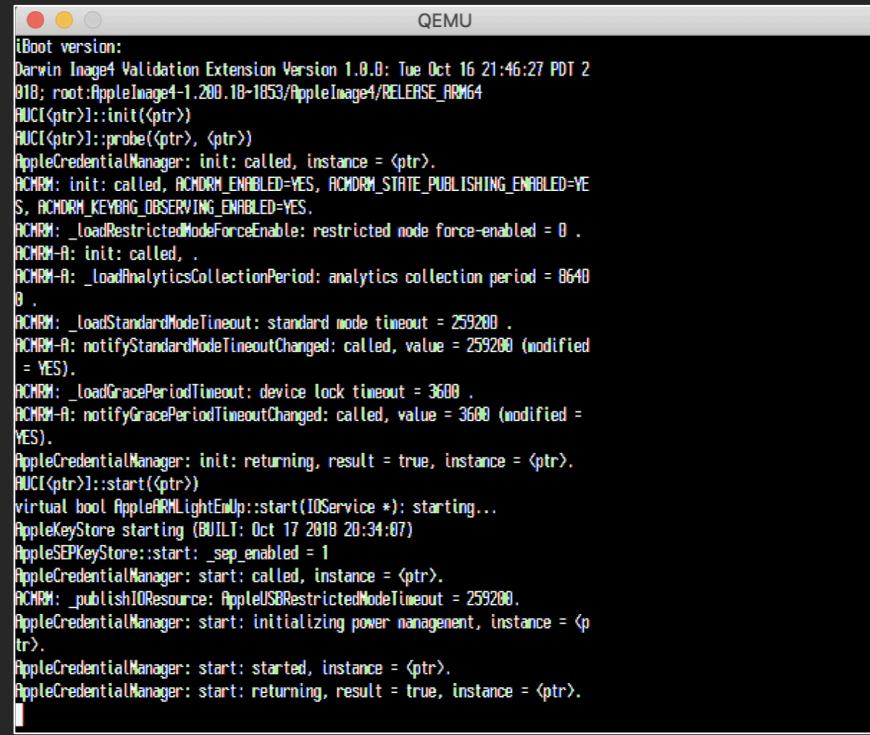
```
#define V_DEPTH      16$  
#define V_HEIGHT     800$  
#define V_WIDTH      600$  
#define V_DISPLAY    0$  
#define V_LINESIZE   (V_WIDTH * 3)$  
-
```

Textual framebuffer

- Connect the QEMU ramfb device fb to the iOS framebuffer

Textual framebuffer

And it works!



The screenshot shows a terminal window titled "QEMU" displaying a textual framebuffer output. The logs include iBoot version information, Darwin Image Validation Extension details, and various system initialization messages from AppleCredentialManager, ACNRM, and AppleKeyStore. The text is white on a black background.

```
iBoot version:  
Darwin Image Validation Extension Version 1.0.0: Tue Oct 16 21:46:27 PDT 2  
018; rootfsAppleImage4-1.200.18-1853/AppleImage4/RELEASE_ARM64  
AUCI(<ptr>):::init(<ptr>)  
AUCI(<ptr>)::probe(<ptr>, <ptr>)  
AppleCredentialManager: init: called, instance = <ptr>.  
ACNRM: init: called, ACNDRM_ENABLED=YES, ACNDRM_STATE_PUBLISHING_ENABLED=YE  
S, ACNDRM_KEYBKG_OBSERVING_ENABLED=YES.  
ACNRM: _loadRestrictedModeForceEnable: restricted mode force-enabled = 0 .  
ACNRM-A: init: called, .  
ACNRM-A: _loadAnalyticsCollectionPeriod: analytics collection period = 8640  
0 .  
ACNRM: _loadStandardModeTimeout: standard mode timeout = 259200 .  
ACNRM-A: notifyStandardModeTimeoutChanged: called, value = 259200 (modified  
= YES).  
ACNRM: _loadGracePeriodTimeout: device lock timeout = 3600 .  
ACNRM-A: notifyGracePeriodTimeoutChanged: called, value = 3600 (modified  
= YES).  
AppleCredentialManager: init: returning, result = true, instance = <ptr>.  
AUCI(<ptr>)::start(<ptr>)  
virtual bool AppleARMLightEndpoint::start(IOService *): starting...  
AppleKeyStore starting (BUILT: Oct 17 2018 20:34:07)  
AppleSEPKeyStore::start: _sep_enabled = 1  
AppleCredentialManager: start: called, instance = <ptr>.  
ACNRM: _publishIOResource: AppleUSBRestrictedModeTimeout = 259200.  
AppleCredentialManager: start: initializing power management, instance = <ptr>.  
AppleCredentialManager: start: started, instance = <ptr>.  
AppleCredentialManager: start: returning, result = true, instance = <ptr>.
```

Aleph Research

Agenda

- Past public research on iOS on QEMU
 - iOS kernel boot process
 - Execution of non-Apple executables with Trust Cache
 - Bash execution with launchd
 - tfp0
 - SSH connection
 - Textual framebuffer
 - **Block device driver**
 - No secure monitor
 - Launchd craziness
 - Next steps
-

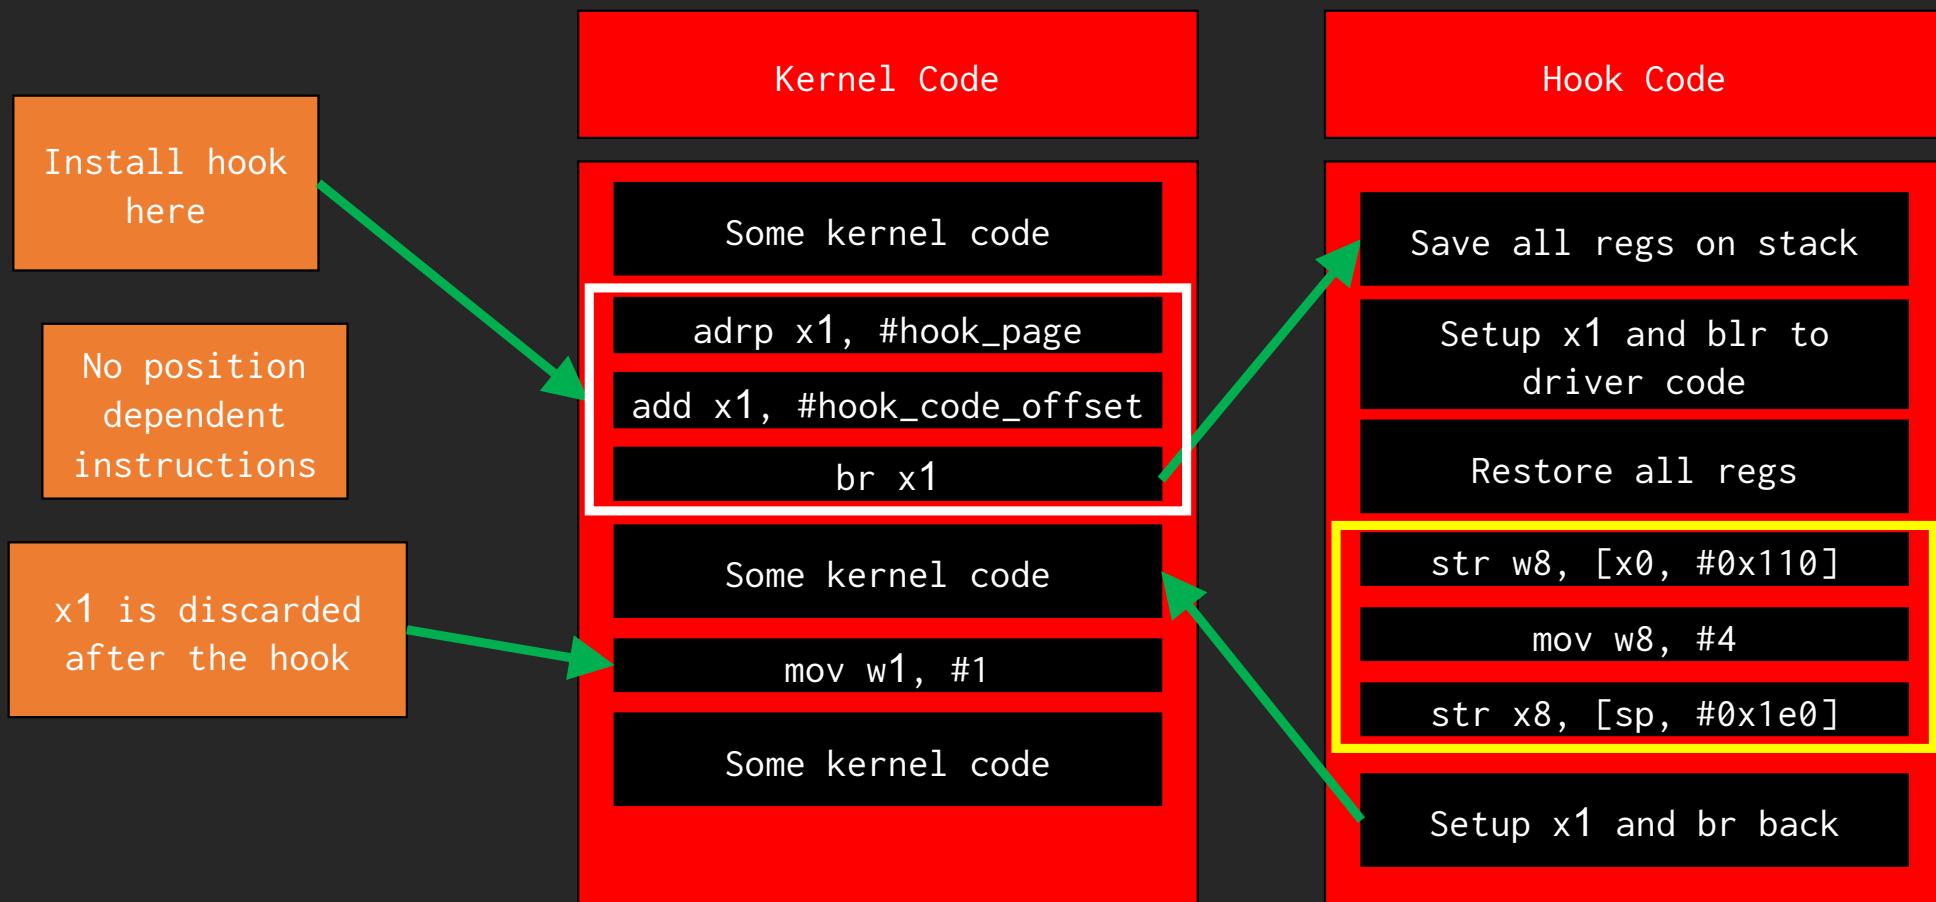
Block Device Driver

- RAMDisk block devices do not support disks larger than 1.5 GB
 - Full disk image is larger than 4GB
 - Using the RAMDisk we only have 1 block device
 - We want a BD for the root mount and another for the R/W data mount - same as a regular iOS system
 - Write our own block device driver to load the 2 large block devices
 - Patch the kernel and call our driver during kernel boot
-

Block Device Driver - Hook

- Allocate kernel memory at a fixed location
- Trampoline hook code in assembly
- Flat binary compilation of c code for the driver
- Parse the page table and make all these pages RWX

Block Device Driver - Hook



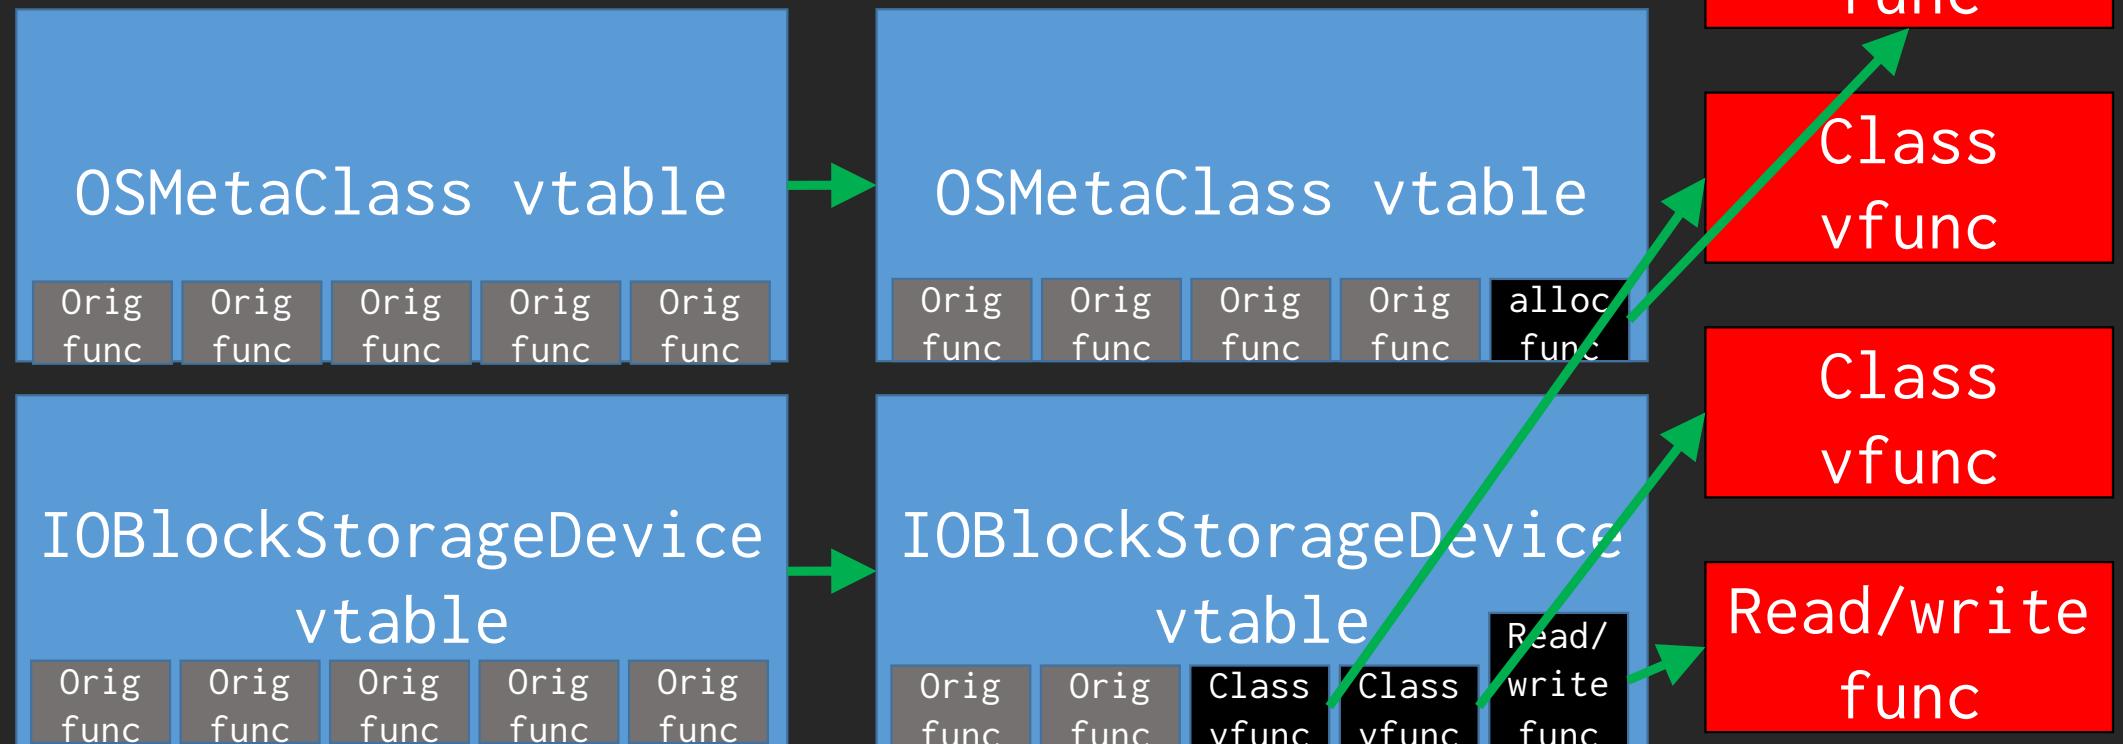
Block Device Driver - Driver init

- We do not have a kernel dev kit for iOS or a compatible cpp compiler
- Create our own driver inheriting from IOBlockStorageDevice with c code and build compatible cpp objects for the driver and its metaclass
- Create vtables for the new class and its OSMetaClass

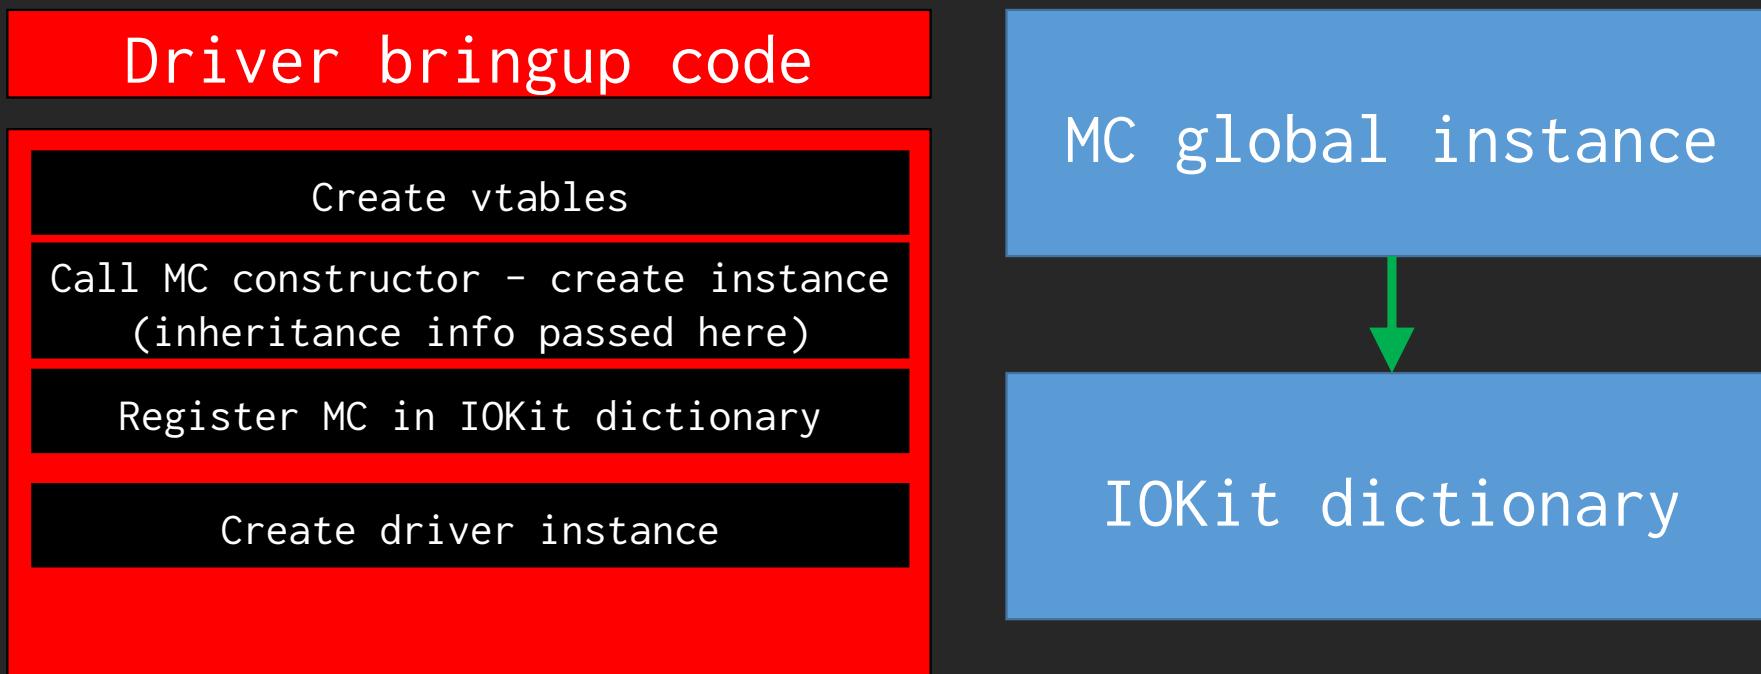
Block Device Driver - Driver init

- Call the OSMetaClass constructor and register the new class in the sAllClassesDict global classes dictionary
- Call `OSMetaClass::allocClassWithName("AlephStorageBlockDevice")` to create 2 instances of our class for the 2 block devices
- Attach them to a node in the IOKit registry and call `::registerService()` on them to enable the new devices

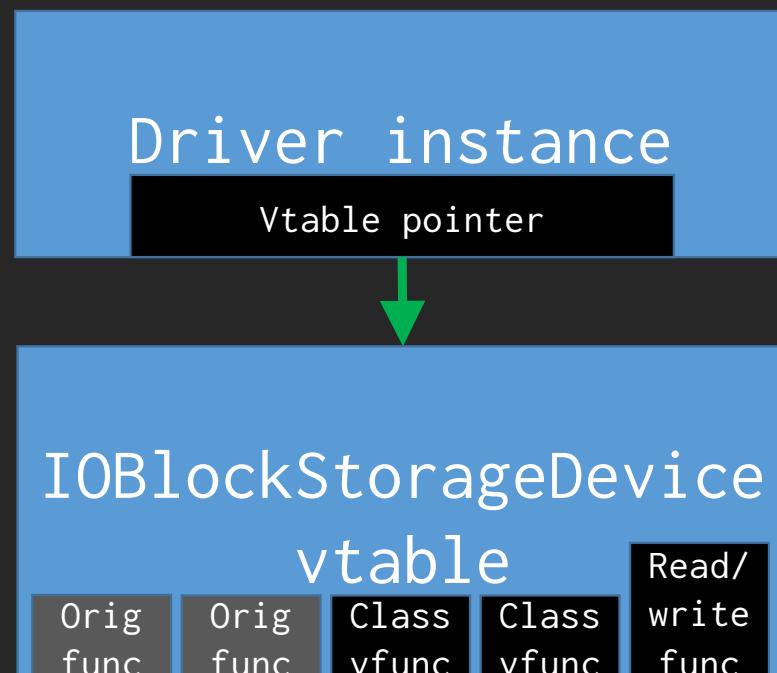
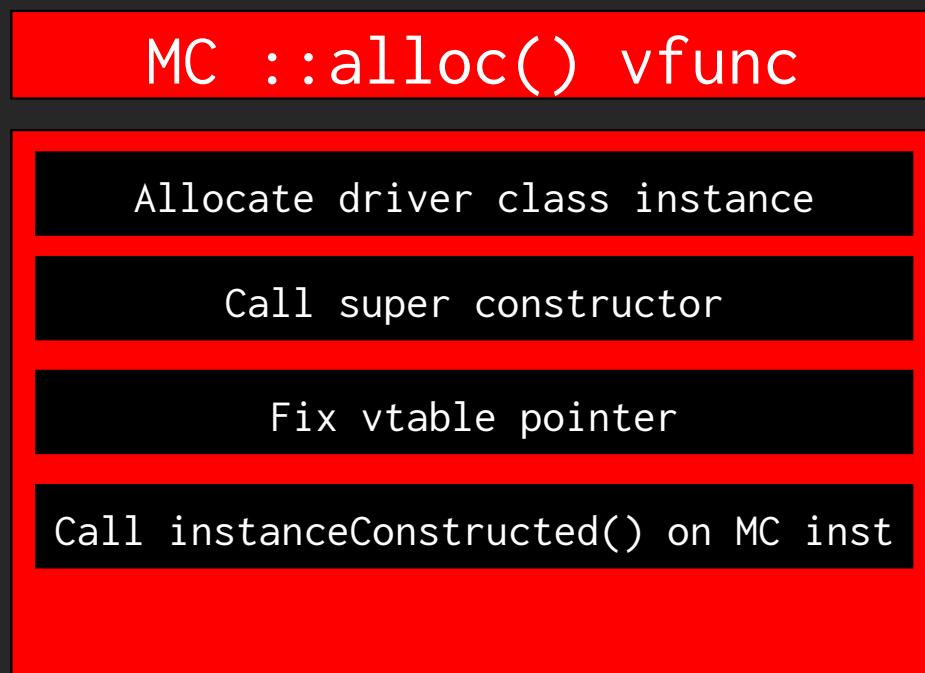
Block Device Driver - Driver init



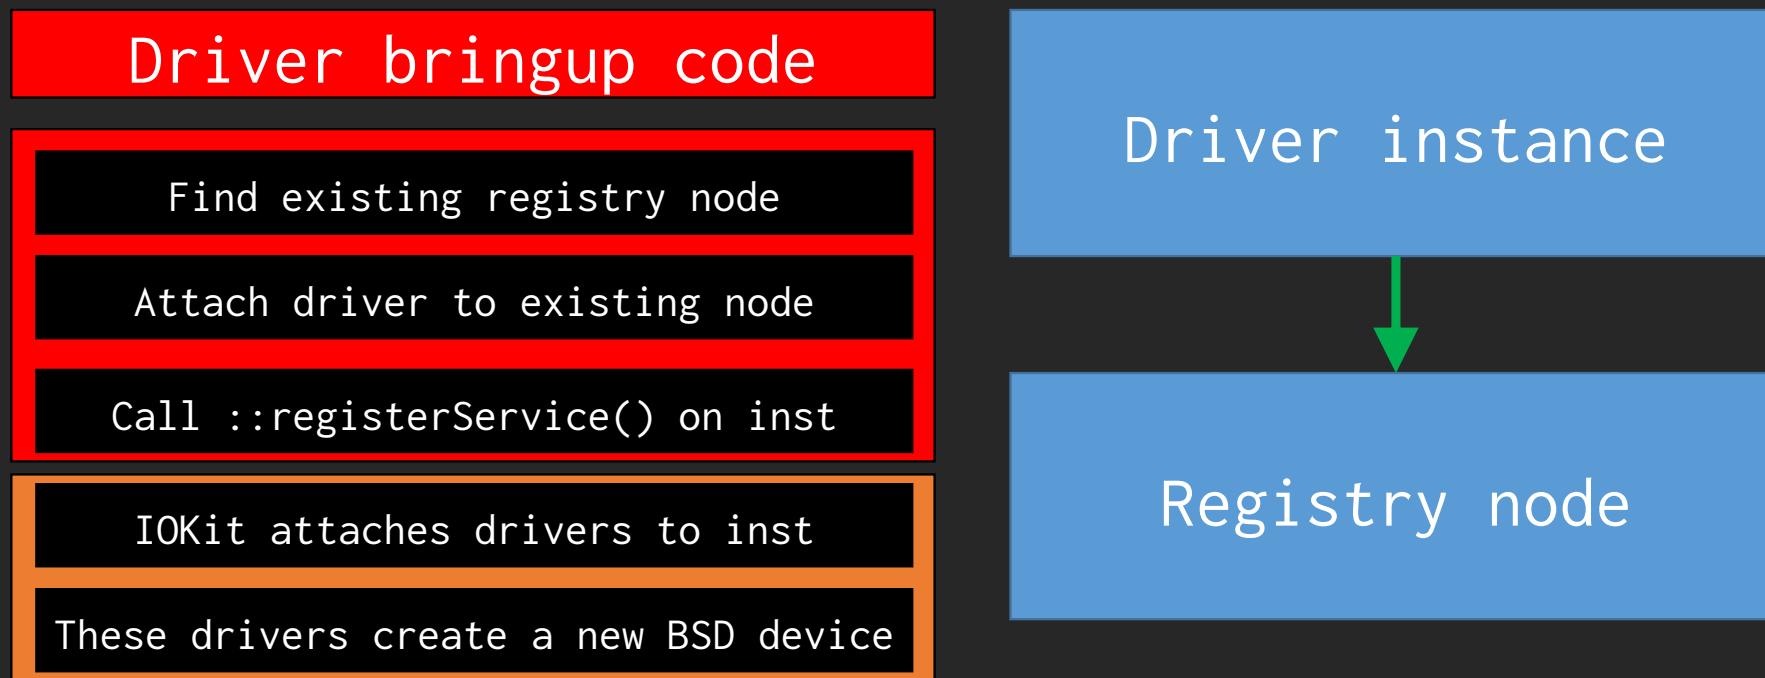
Block Device Driver - Driver init



Block Device Driver - Driver init



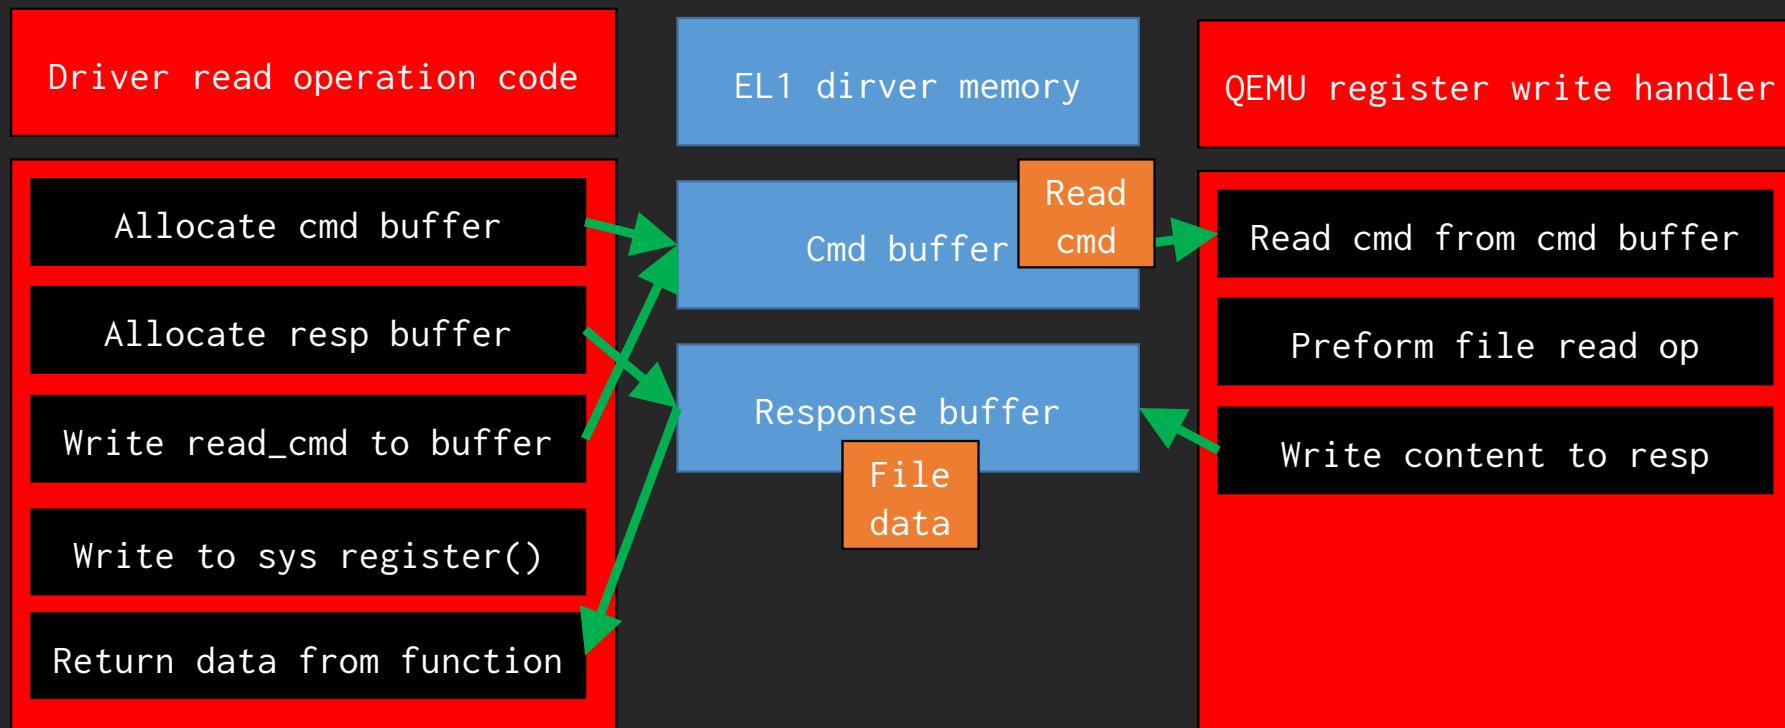
Block Device Driver - Driver init



Block Device Driver - Driver read/write

- Implement the virtual functions to read/write from/to the new block devices
- Write to a system register and the QEMU callback reads/writes directly from/to memory buffers on the guest to transfer the data
- The block device data on the QEMU host is backed by files

Block Device Driver - Driver read/write



```
+o AlephStorageBlockDevice <class IORegistryEntry:IOService:IOBlockStorageDevice:AlephStorageBlockDevice, id 0x1000001c7, registered, matched, active, busy 0 (633 ms), retain 7>
{
    "device-type" = "Generic"
    "IOMinimumSegmentAlignmentByteCount" = 0x4
}

+-o IOBlockStorageDriver <class IORegistryEntry:IOService:IOStorage:IOBlockStorageDriver, id 0x1000001c8, registered, matched, active, busy 0 (613 ms), retain 8>
{
    "IOPropertyMatch" = {"device-type"="Generic"}
    "IOProbeScore" = 0x0
    "IOProviderClass" = "IOBlockStorageDevice"
    "IOClass" = "IOBlockStorageDriver"
    "CFBundleIdentifier" = "com.apple.iokit.IOStorageFamily"
    "Statistics" = {"Operations (Write)"=0x0,"Latency Time (Write)"=0x0,"Bytes (Read)"=0x1a5f9000,"Errors (Write)"=0x0,"Total Time (Read)"=0x427c05e97,"Retries (Read)"=0x0,"Latencies (Read)"=0x13cb,"Retries (Write)"=0x0}
    "IOMatchCategory" = "IODefaultMatchCategory"
    "IOGeneralInterest" = "IOCommand is not serializable"
}

+-o 0Aleph 0AlephBDev Media <class IORegistryEntry:IOService:IOStorage:IOMedia, id 0x1000001c9, registered, matched, active, busy 0 (613 ms), retain 9>
{
    "Removable" = No
    "Content" = ""
    "Whole" = Yes
    "Leaf" = Yes
    "BSD Name" = "disk0"
    "Ejectable" = No
    "Preferred Block Size" = 0x1000
    "IOMediaIcon" = {"IOBundleResourceFile"="External.icns","CFBundleIdentifier"="com.apple.iokit.IOStorageFamily"}
    "BSD Minor" = 0x0
    "Writable" = Yes
    "BSD Major" = 0x1
    "Size" = 0x180000000
    "Open" = Yes
    "Content Hint" = ""
    "BSD Unit" = 0x0
}

+-o IOMediaBSDClient <class IORegistryEntry:IOService:IOMediaBSDClient, id 0x1000001ca, registered, matched, active, busy 0 (3 ms), retain 7>
{
    "IOClass" = "IOMediaBSDClient"
    "IOMatchCategory" = "IOMediaBSDClient"
    "IOProbeScore" = 0x7530
    "IOProviderClass" = "IOMedia"
    "IOResourceMatch" = "IOBSD"
    "CFBundleIdentifier" = "com.apple.iokit.IOStorageFamily"
}
```

Block Device Driver

```
bash-4.4# mount -t hfs /dev/disk1 /mnt1

mount_hfs: Could not create property for re-key environment check: No such file or directory
hfs: mounted PeaceB16B92.arm64UpdateRamDisk on device disk1
bash-4.4#
bash-4.4# cd /mnt1/
bash-4.4# ls
.fseventsda      buddy          log             run
Keychains        containers     logs            staged_system_apps
Managed Preferences db            mobile           tmp
MobileAsset       dropbear       msgs            vm
MobileDevice     empty          networkd        wireless
MobileSoftwareUpdate folders        preferences
audit            installd       root
bash-4.4#
```

Agenda

- Past public research on iOS on QEMU
 - iOS kernel boot process
 - Execution of non-Apple executables with Trust Cache
 - Bash execution with launchd
 - tfp0
 - SSH connection
 - Textual framebuffer
 - Block device driver
 - **No secure monitor**
 - Launchd craziness
 - Next steps
-

No Secure Monitor

- Move from not patching the kernel to patching it for driver loading
- Disable the Secure Monitor to disable KPP

No Secure Monitor

- Don't load the Secure Monitor image
- Start executing the kernel at EL1 at its entry point
- NOP SMC instruction (Secure Monitor Call)

Register group: general											
x0	0x800	2048	x1	0x470a5000	1191858176	x2	0x0	0			
x6	0xff	255	x7	0xad0	2768	x8	0x40000000	1073741824			
x12	0x14	8	x13	0x10000000	16777216	x14	0x3000000	50331648	[32]		
x18	0x0	0	x19	0xfffffff75d1000	-68595937280	x20	0xfffffff0070a5088	-68601360248			
x24	0xf420	1000000	x25	0xfffffff00707d000	-68601524224	x26	0xfffffff007078440	-68601543616			
x30	0xfffffffff0071b74a0	-68600236896	sp	0xfffffff00767fe10	0xfffffff0070a7d3c		0xfffffff0070a7d3d	0xfffffff0070a7d3d			
MVFR6_EL1_RESERVED	0x0	0	MVFR4_EL1_RESERVED	0x0	0	ID_AA64PFR1_EL1	0x0	0			
ID_AA64PFR5_EL1_RESERVED	0x0	0	ID_AA64PFR6_EL1_RESERVED	0x0	0	ID_AA64PFR7_EL1_RESERVED	0x0	0			
ID_AA64DFR3_EL1_RESERVED	0x0	0	ID_AA64AFR1_EL1	0x0	0	ID_AA64AFR2_EL1_RESERVED	0x0	0			
ID_AA64ISARS_EL1_RESERVED	0x0	0	ID_AA64ISAR3_EL1_RESERVED	0x0	0	ID_AA64ISAR1_EL1	0x0	0			
ID_AA64MMFR0_EL1	0x0	0	ID_AA64MMFR1_EL1	0x0	0	ID_AA64MMFR2_EL1_RESERVED	0x0	0			
ID_AA64MMFR0_EL1_RESERVED	0x0	0	ID_AA64MMFR0_EL1	0x124	4388	ID_AA64MMFR5_EL1_RESERVED	0x0	0			
REVIDR_EL1	0x0	0	SCTRLR	0x3454593d	877943101	ACTLR_EL1	0x0	0			
ISCTRLR_EL2	0x0	0	HSTR_EL2	0x0	0	CPTLR_EL2	0x0	0			
MDCR_EL2	0x0	0	MDCR_EL3	0x0	0	SCTRLR_EL3	0x30d5180d	819271693			
PMCR_EL0	0x41000000	1090519840	PMCNTENCLR_EL0	0x0	0	PMCNTESTEL0	0x0	0			
PMCCNTREL0	0x0	0	PMCEID1_EL0	0x0	0	TTBR0_EL1	0x1000007a40000	281477789253632			
MAIR_EL1	0x44f00bb44ff	4737361200383	TTBR0_EL2	0x0	0	TCR_EL2	0x0	0			
TTBR0_EL3	0x410000c000	27917293392	MAIR_EL2	0x0	0	TCR_EL3	0x1a511	107793			
L2CTLR_EL1	0x0	0	L2CTLR_EL1	0x0	0	DACK32_EL2	0x0	0			
SP_EL0	0x0	0	VBAR_EL2	0x0	0	SP_EL1	0xfffffff007688000	-68595187712			
SPSR_IRO	0x0	0	SPSR_ABТ	0x0	0	FPCR	0x0	0			
0x4fffff0070a7d3c	smc #0x11										
0xfffffff0070a7d40	ret										
0xfffffff0070a7d44	.inst 0x00000000	; undefined									
0xfffffff0070a7d45	.inst 0x00000000	; undefined									
0xfffffff0070a7d46	.inst 0x00000000	; undefined									
0xfffffff0070a7d47	.inst 0x00000000	; undefined									
0xfffffff0070a7d48	.inst 0x00000000	; undefined									
0xfffffff0070a7d49	.inst 0x00000000	; undefined									
0xfffffff0070a7d4a	.inst 0x00000000	; undefined									
0xfffffff0070a7d4b	.inst 0x00000000	; undefined									
0xfffffff0070a7d4c	.inst 0x00000000	; undefined									
0xfffffff0070a7d4d	.inst 0x00000000	; undefined									
0xfffffff0070a7d4e	.inst 0x00000000	; undefined									
0xfffffff0070a7d4f	.inst 0x00000000	; undefined									
0xfffffff0070a7d50	.inst 0x00000000	; undefined									
0xfffffff0070a7d51	.inst 0x00000000	; undefined									
0xfffffff0070a7d52	.inst 0x00000000	; undefined									
0xfffffff0070a7d53	.inst 0x00000000	; undefined									
0xfffffff0070a7d54	.inst 0x00000000	; undefined									
0xfffffff0070a7d55	.inst 0x00000000	; undefined									
0xfffffff0070a7d56	.inst 0x00000000	; undefined									
0xfffffff0070a7d57	.inst 0x00000000	; undefined									
0xfffffff0070a7d58	.inst 0x00000000	; undefined									
0xfffffff0070a7d59	.inst 0x00000000	; undefined									
0xfffffff0070a7d5a	.inst 0x00000000	; undefined									
0xfffffff0070a7d5b	.inst 0x00000000	; undefined									
0xfffffff0070a7d5c	.inst 0x00000000	; undefined									
0xfffffff0070a7d5d	.inst 0x00000000	; undefined									
0xfffffff0070a7d5e	.inst 0x00000000	; undefined									
0xfffffff0070a7d5f	.inst 0x00000000	; undefined									
0xfffffff0070a7d60	.inst 0x00000000	; undefined									
0xfffffff0070a7d61	.inst 0x00000000	; undefined									
0xfffffff0070a7d62	.inst 0x00000000	; undefined									
0xfffffff0070a7d63	.inst 0x00000000	; undefined									
0xfffffff0070a7d64	.inst 0x00000000	; undefined									
0xfffffff0070a7d65	.inst 0x00000000	; undefined									
0xfffffff0070a7d66	.inst 0x00000000	; undefined									
0xfffffff0070a7d67	.inst 0x00000000	; undefined									
0xfffffff0070a7d68	.inst 0x00000000	; undefined									
0xfffffff0070a7d69	.inst 0x00000000	; undefined									
0xfffffff0070a7d6a	.inst 0x00000000	; undefined									
0xfffffff0070a7d6b	.inst 0x00000000	; undefined									
0xfffffff0070a7d6c	.inst 0x00000000	; undefined									
0xfffffff0070a7d6d	.inst 0x00000000	; undefined									
0xfffffff0070a7d6e	.inst 0x00000000	; undefined									
0xfffffff0070a7d6f	.inst 0x00000000	; undefined									
0xfffffff0070a7d70	.inst 0x00000000	; undefined									
0xfffffff0070a7d71	.inst 0x00000000	; undefined									
0xfffffff0070a7d72	.inst 0x00000000	; undefined									
0xfffffff0070a7d73	.inst 0x00000000	; undefined									
0xfffffff0070a7d74	.inst 0x00000000	; undefined									
0xfffffff0070a7d75	.inst 0x00000000	; undefined									
0xfffffff0070a7d76	.inst 0x00000000	; undefined									
0xfffffff0070a7d77	.inst 0x00000000	; undefined									
0xfffffff0070a7d78	.inst 0x00000000	; undefined									
0xfffffff0070a7d79	.inst 0x00000000	; undefined									
0xfffffff0070a7d7a	.inst 0x00000000	; undefined									
0xfffffff0070a7d7b	.inst 0x00000000	; undefined									
0xfffffff0070a7d7c	.inst 0x00000000	; undefined									
0xfffffff0070a7d7d	.inst 0x00000000	; undefined									
0xfffffff0070a7d7e	.inst 0x00000000	; undefined									
0xfffffff0070a7d7f	.inst 0x00000000	; undefined									
0xfffffff0070a7d80	.inst 0x00000000	; undefined									
0xfffffff0070a7d81	.inst 0x00000000	; undefined									
0xfffffff0070a7d82	.inst 0x00000000	; undefined									
0xfffffff0070a7d83	.inst 0x00000000	; undefined									
0xfffffff0070a7d84	.inst 0x00000000	; undefined									
0xfffffff0070a7d85	.inst 0x00000000	; undefined									
0xfffffff0070a7d86	.inst 0x00000000	; undefined									
0xfffffff0070a7d87	.inst 0x00000000	; undefined									
0xfffffff0070a7d88	.inst 0x00000000	; undefined									
0xfffffff0070a7d89	.inst 0x00000000	; undefined									
0xfffffff0070a7d8a	.inst 0x00000000	; undefined									
0xfffffff0070a7d8b	.inst 0x00000000	; undefined									
0xfffffff0070a7d8c	.inst 0x00000000	; undefined									
0xfffffff0070a7d8d	.inst 0x00000000	; undefined									
0xfffffff0070a7d8e	.inst 0x00000000	; undefined									
0xfffffff0070a7d8f	.inst 0x00000000	; undefined									
0xfffffff0070a7d90	.inst 0x00000000	; undefined									
0xfffffff0070a7d91	.inst 0x00000000	; undefined									
0xfffffff0070a7d92	.inst 0x00000000	; undefined									
0xfffffff0070a7d93	.inst 0x00000000	; undefined									
0xfffffff0070a7d94	.inst 0x00000000	; undefined									
0xfffffff0070a7d95	.inst 0x00000000	; undefined									
0xfffffff0070a7d96	.inst 0x00000000	; undefined									
0xfffffff0070a7d97	.inst 0x00000000	; undefined									
0xfffffff0070a7d98	.inst 0x00000000	; undefined									
0xfffffff0070a7d99	.inst 0x00000000	; undefined									
0xfffffff0070a7d9a	.inst 0x00000000	; undefined									
0xfffffff0070a7d9b	.inst 0x00000000	; undefined									
0xfffffff0070a7d9c	.inst 0x00000000	; undefined									
0xfffffff0070a7d9d	.inst 0x00000000	; undefined									
0xfffffff0070a7d9e	.inst 0x00000000	; undefined									
0xfffffff0070a7d9f	.inst 0x00000000	; undefined									
0xfffffff0070a7d99	.inst 0x00000000	; undefined									
0xfffffff0070a7d9a	.inst 0x00000000	; undefined									
0xfffffff0070a7d9b	.inst 0x00000000	; undefined									
0xfffffff0070a7d9c	.inst 0x00000000	; undefined									
0xfffffff0070a7d9d	.inst 0x00000000	; undefined									
0xfffffff0070a7d9e	.inst 0x00000000	; undefined									
0xfffffff0070a7d9f	.inst 0x00000000	; undefined									
0xfffffff0070a7d99	.inst 0x00000000	; undefined									
0xfffffff0070a7d9a	.inst 0x00000000	; undefined									
0xfffffff0070a7d9b	.inst 0x00000000	; undefined									
0xfffffff0070a7d9c	.inst 0x00000000	; undefined									
0xfffffff0070a7d9d	.inst 0x00000000	; undefined									
0xfffffff0070a7d9e	.inst 0x00000000	; undefined									
0xfffffff0070a7d9f	.inst 0x00000000	; undefined									
0xfffffff0070a7d99	.inst 0x00000000	; undefined									
0xfffffff0070a7d9a	.inst 0x00000000	; undefined		</							

No Secure Monitor

- Another way KPP kicks in is by trapping FP operations done in EL0 or in EL1 to EL3 (secure monitor)
- Intercept writes to CPACR_EL1 (in QEMU) so that it will always hold 0x30000
 - This means FPEN is 3 (don't trap any FP execution at all)
- [@xerub](https://xerub.github.io/ios/kpp/2017/04/13/tick-tock.html)

Agenda

- Past public research on iOS on QEMU
 - iOS kernel boot process
 - Execution of non-Apple executables with Trust Cache
 - Bash execution with launchd
 - tfp0
 - SSH connection
 - Textual framebuffer
 - Block device driver
 - No secure monitor
 - **Launchd craziness**
 - Next steps
-

Launchd craziness

- Created 2 raw HFS block device disk images with the full disk images content from the software update package
- Started QEMU pointing the root mount to the new block devices
- Added launchd items to execute our bash shell, start the TCP tunnel, start the dropbear SSH server and other services
- Our services didn't run but the original system services did run

Launchd craziness

Any ideas?

Launchd craziness

- Removed all launchd items
- Still no go - they execute the same as before
- Removed the directory
- Panic and nothing executes
- Boot with the content of the ramdisk we used before
- Panic with an error about a missing service cache

Launchd craziness

```
 iVar2 = _dlopen("/System/Library/Caches/com.apple.xpcd/xpcd_cache.dylib",2);
if ((iVar2 != 0) &&
    (UNRECOVERED_JUMPTABLE = (code *)_dlsym(iVar2,"__xpcd_cache"),
     UNRECOVERED_JUMPTABLE != (code *)0x0)) {
    iVar1 = (*UNRECOVERED_JUMPTABLE)();
    if (iVar1 != 1) {
        __os_assert_log(0);
        __os_crash();
        /* WARNING: Could not recover jumptable at 0x00010000a248. Too many branches */
        /* WARNING: Treating indirect jump as call */
        UNRECOVERED_JUMPTABLE = (code *)SoftwareBreakpoint(1,0x10000a24c);
        (*UNRECOVERED_JUMPTABLE)();
        return;
    }
    iVar1 = _dladdr(UNRECOVERED_JUMPTABLE,auStack72);
    if (iVar1 != 0) {
        puVar3 = _getsectiondata(local_40,"__TEXT",__xpcd_cache,(ulong *)&local_28);
        if (puVar3 == (uint8_t *)0x0) {
            DAT_1000436b8 = _xpc_dictionary_create(0,0,0);
        }
        else {
            DAT_1000436b8 = _xpc_create_from_plist(puVar3,local_28);
        }
    }
    if ((DAT_1000436b8 == 0) ||
        (puVar4 = (undefined *)_xpc_get_type(), puVar4 != __xpc_type_dictionary)) {
        if (DAT_100044479 == '\0') {
            /* WARNING: Subroutine does not return */
            FUN_100028f84("No service cache");
        }
        DAT_1000436b8 = _xpc_dictionary_create(0,0,0);
    }
    DAT_100043670 = _xpc_dictionary_get_value(DAT_1000436b8,"LaunchDaemons");
    if ((DAT_100043670 == 0) ||
        (puVar4 = (undefined *)_xpc_get_type(), puVar4 != __xpc_type_dictionary)) {
        if (DAT_100044479 == '\0') {
            /* WARNING: Subroutine does not return */
            FUN_100028f84("No daemons in cache");
        }
    }
}
```

Launchd craziness

- Launchd decides whether to use this cache or the filesystem based on the bootargs (if it thinks it is a ramdisk boot)
- Patch this code “finds” rd=md0 in the boot args

```
lVar5 = FUN_100029448("kern.bootargs",&local_48);
if (lVar5 == 0) {
    cVar10 = '\0';
}
else {
    pcVar6 = _strnstr(local_48,"rd=md0",(size_t)lVar5);
    cVar10 = (char)pcVar6 + '\x01';
}
_free(local_48);
DAT_100044479 = cVar10;
```

Launchd craziness

- Sign the patched binary
- Update the static trust cache
- Replace the launchd binary in the filesystem
- Remove the service cache dynamic lib file from the filesystem

Launchd craziness

And it works!

Launchd craziness

*	[0]	"kernel_task"	Regular	12	0xffffffe0005f9c20
	[1]	"launchd"			0xffffffe0005fa760
	[19]	"xpcproxy"			0xffffffe000c305a0
0	[27]	"xpcproxy" 4		6	0xffffffe000c61680
	[39]	"xpcproxy"			0xffffffe000c98b40
	[44]	"xpcproxy"			0xffffffe000cbd680
	[52]	"xpcproxy"			0xffffffe000cbc5a0
	[60]	"xpcproxy"			0xffffffe000cf1c0
	[62]	"xpcproxy"			0xffffffe000cfed00
	[6]	"tseventsd"			0xffffffe000cff2a0
	[7]	"bash"			0xffffffe000cff840
	[16]	"configd"			0xffffffe0005fb2a0
	[36]	"SpringBoard"			0xffffffe000d4d680
	[21]	"tunnel"			0xffffffe000d4cb40
	[23]	"wiflu"			0xffffffe000d4dc20
	[28]	"mobiletimerd"			0xffffffe000d4e1c0
	[19]	"xpcproxy"			0xffffffe000d4ed00
	[62]	"xpcproxy"			0xffffffe000d4f2a0
	[45]	"locationd"			0xffffffe000d4f840
	[42]	"dropbear"			0xffffffe000d4c000
	[60]	"xpcproxy"			0xffffffe000c33840
	[39]	"xpcproxy"			0xffffffe000c310e0
	[55]	"AppleCredentialManagerDaemon"			0xffffffe000c332a0
	[43]	"timed"			0xffffffe000c990e0
	[50]	"mobilewatchdog"			0xffffffe000c99680
	[51]	"ptpd"			0xffffffe0005f9680
	[52]	"xpcproxy"			0xffffffe0005f8b40
	[59]	"nanopoold"			0xffffffe000c62760

```
vera — vim /Users/vera/OpenSourceGitProjects/darwin-xnu/osfmk/kern/sched_prim.c — vim — vim ~/OpenSourceGitProjects/darwin-xnu/osfmk/kern/sched_prim.c — 103x26

static boolean_t
thread_invoke(
    thread_t self,
    thread_t thread,
    ast_t reason)

    if (__improbable(get_preemption_level() != 0)) {
        int pl = get_preemption_level();
        panic("thread_invoke: preemption_level %d, possible cause: %s",
              pl, (pl < 0 ? "unlocking an unlocked mutex or spinlock" :
                    "blocking while holding a spinlock, or within interrupt context"));
    }

    thread_continue_t continuation = self->continuation;
    void *parameter = self->parameter;
    processor_t processor;

    uint64_t ctime = mach_absolute_time();

#endif CONFIG_MACH_APPROXIMATE_TIME
    commpage_update_mach_approximate_time(ctime);
#endif

#if defined(CONFIG_SCHED_TIMESHARE_CORE)
    if ((thread->state & TH_IDLE) == 0)
```

Demo - Research a vulnerability - voucher_swap

- Research done by Brandon Azad
- iOS 12.1 jailbreak
- Trigger the vulnerability while debugging

```
1) ios_command_line_tool.m
}$
$
int main(int argc, const char *argv[]) {
    kern_return_t kr;
    mach_port_t thread;
    mach_port_t uaf_port;
    mach_port_t discloser_mach_port = MACH_PORT_NULL;

    printf("start PoC\n");

    kr = thread_create(mach_task_self(), &thread);

    uaf_port = create_voucher(0);

    kr = thread_set_mach_voucher(thread, uaf_port);

    voucher_release(uaf_port);

    mach_port_destroy(mach_task_self(), uaf_port);

    thread_get_mach_voucher(thread, 0, &discloser_mach_port);
    return 0;
}
$
$
$
$
$
$
$

ios_command_line_tool.m:
```

106-1 100% 00

[1 bash] [2 bash] [3 bash] [4 bash] [5 bash] [6 bash] [7 bash]

Sunday 17:29 17/11/2019

Agenda

- Past public research on iOS on QEMU
 - iOS kernel boot process
 - Execution of non-Apple executables with Trust Cache
 - Bash execution with launchd
 - tfp0
 - SSH connection
 - Textual framebuffer
 - Block device driver
 - No secure monitor
 - Launchd craziness
 - **Next steps**
-

Next steps and challenges

- GUI and HID support
 - SEP and keybagd (if required for regular functioning)
 - Interrupt support for TCP tunnel performance improvements
 - VPN connection testing over SSH
 - Performance improvements
 - Virtualization instead of emulation (qemu-kvm)
 - More than 1 CPU
 - Understand what it means to use a real GPU for OpenGL ES
 - More iOS versions and iOS devices
 - Work on symbolication with Jonathan Levin's jtool/joker and otherwise for more versions
 - Fast state recovery
 - Security research and fast coverage-guided, structure-aware with state recovery fuzzing
-

HLT

- Use the the project and contribute!
<https://github.com/alephsecurity>
- Check out our blog: <https://alephsecurity.com>
- Follow us on twitter: @alephsecurity @JonathanAfek
- Questions?

