

Strategic Sponsors



Gold Sponsors



Silver Sponsors





Sasha Goldshtein, CTO, Sela Group
Microsoft Regional Director, C# MVP
@goldshn

Debugging and Profiling .NET Core Apps on Linux

The Plan

This is a talk on debugging and profiling .NET Core apps on Linux—yes, it's pretty crazy that we got this far!

You'll learn:

- ☐ To profile CPU activity in .NET Core apps
- ☐ To visualize stack traces (e.g. of CPU samples) using flame graphs
- ☐ To use Linux tracing tools with .NET Core processes
- ☐ To capture .NET Core runtime events using LTTng
- ☐ To generate and analyze core dumps of .NET Core apps

Disclaimer

Some of this stuff is changing quite rapidly

The tools described here sort of work for .NET Core 2.0, but your mileage may vary

Some of this relies on scripts I hacked together, and will hopefully be officially supported in the future

Tools And Operating Systems Supported



	Linux	Windows	macOS
CPU sampling	perf, BCC	ETW	Instruments, dtrace
Dynamic tracing	perf, SystemTap, BCC	⊖	dtrace
Static tracing	LTTng	ETW	⊖
Dump generation	core_pattern, gcore	Procdump, WER	kern.corefile, gcore
Dump analysis	lldb	Visual Studio, WinDbg	lldb
	This talk		

⚠ Mind The Overhead

Any observation can change the state of the system, but some observations are worse than others

Diagnostic tools have overhead

- Check the docs

- Try on a test system first

- Measure degradation introduced by the tool

OVERHEAD


This traces various kernel page cache functions and maintains in-kernel counts, which are asynchronously copied to user-space. While the rate of operations can be very high (>1G/sec) we can have up to 34% overhead, this is still a relatively efficient way to trace these events, and so the overhead is expected to be small for normal workloads. Measure in a test environment.

—*man cachestat (from BCC)*


Sampling vs. Tracing

Sampling works by getting a snapshot or a call stack every N occurrences of an interesting event

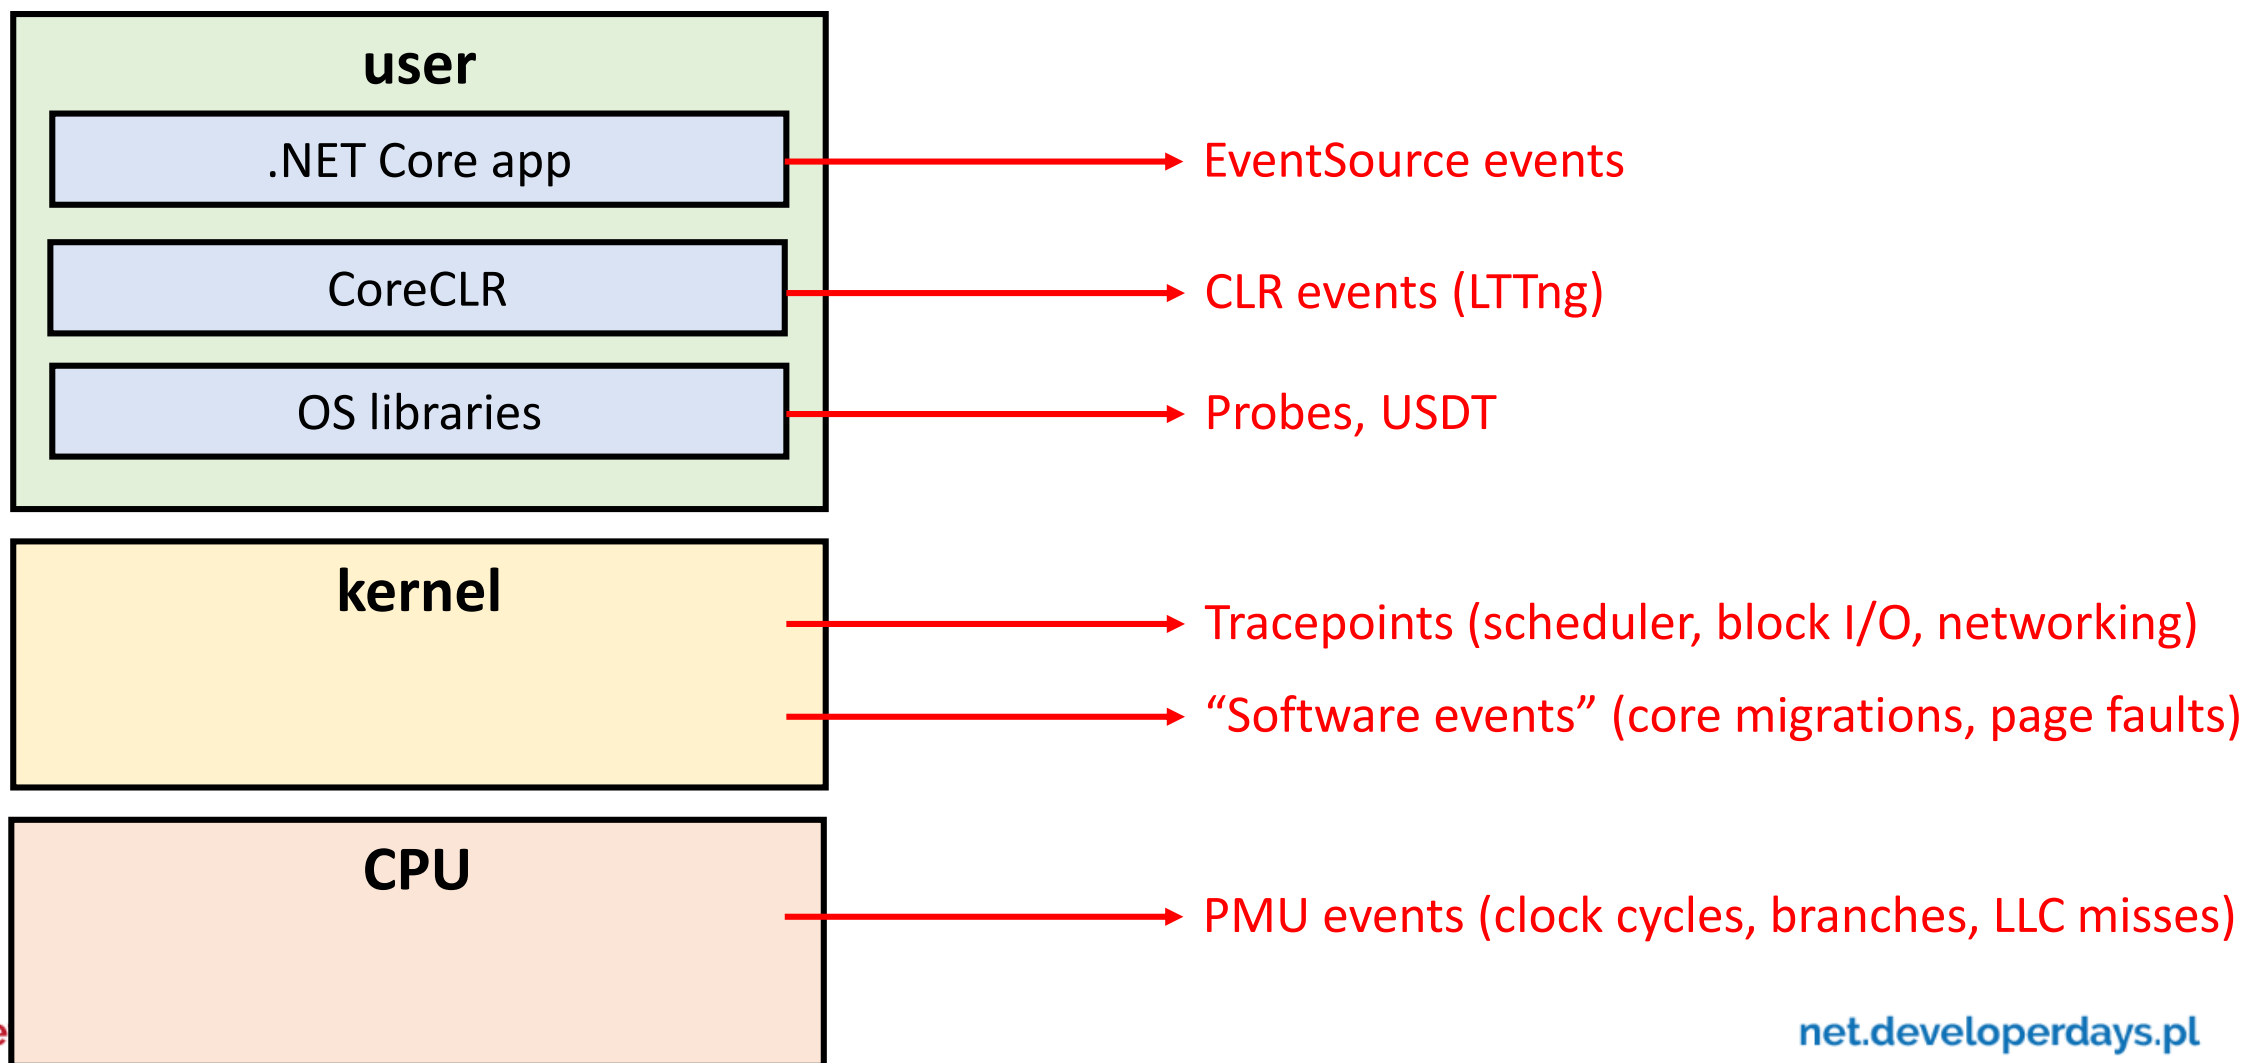
For most events, implemented in the PMU using overflow counters and interrupts



Tracing works by getting a message or a call stack at every occurrence of an interesting event



.NET Core on Linux Tracing Architecture



The Official Story: **perfcollect** and PerfView

1. Download [perfcollect](#)
2. Install prerequisites: `./perfcollect install`
3. Run collection: `./perfcollect collect mytrace`
4. Copy the mytrace.zip file to a Windows machine 🤖
5. Download [PerfView](#) 😁
6. Open the trace in PerfView 🙌

perf

perf is a Linux multi-tool for performance investigations

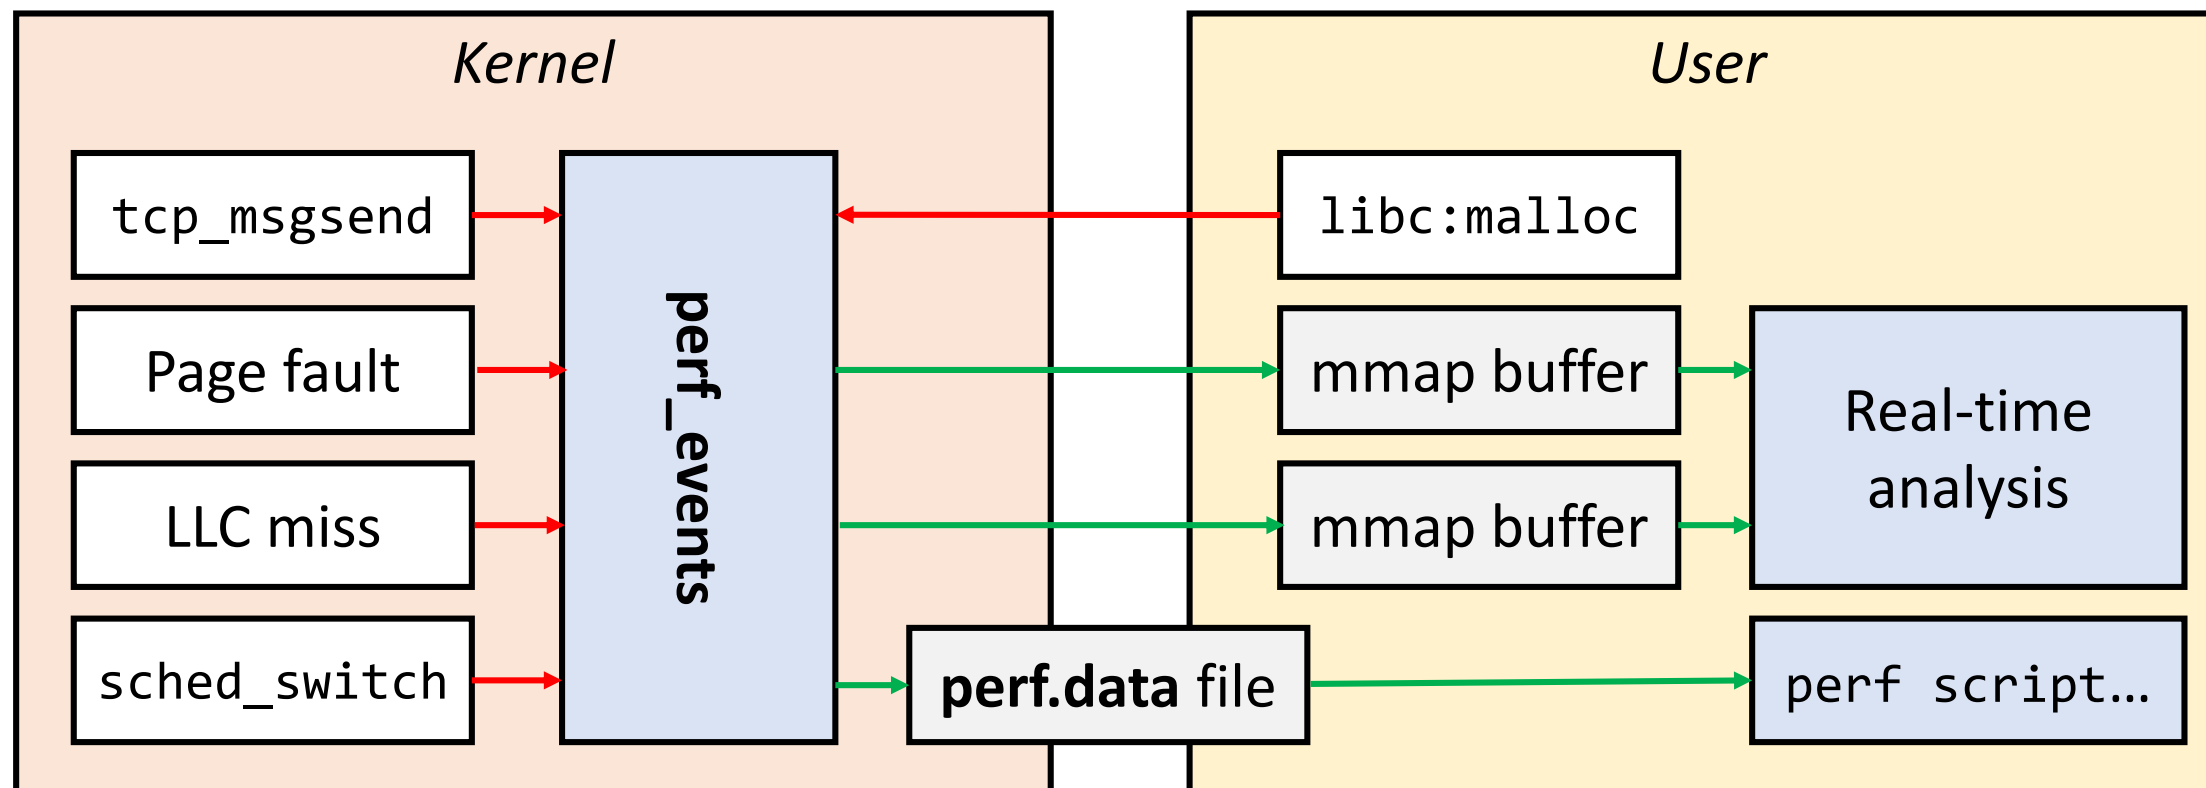
Capable of both tracing and sampling

Developed in the kernel tree, must match running kernel's version

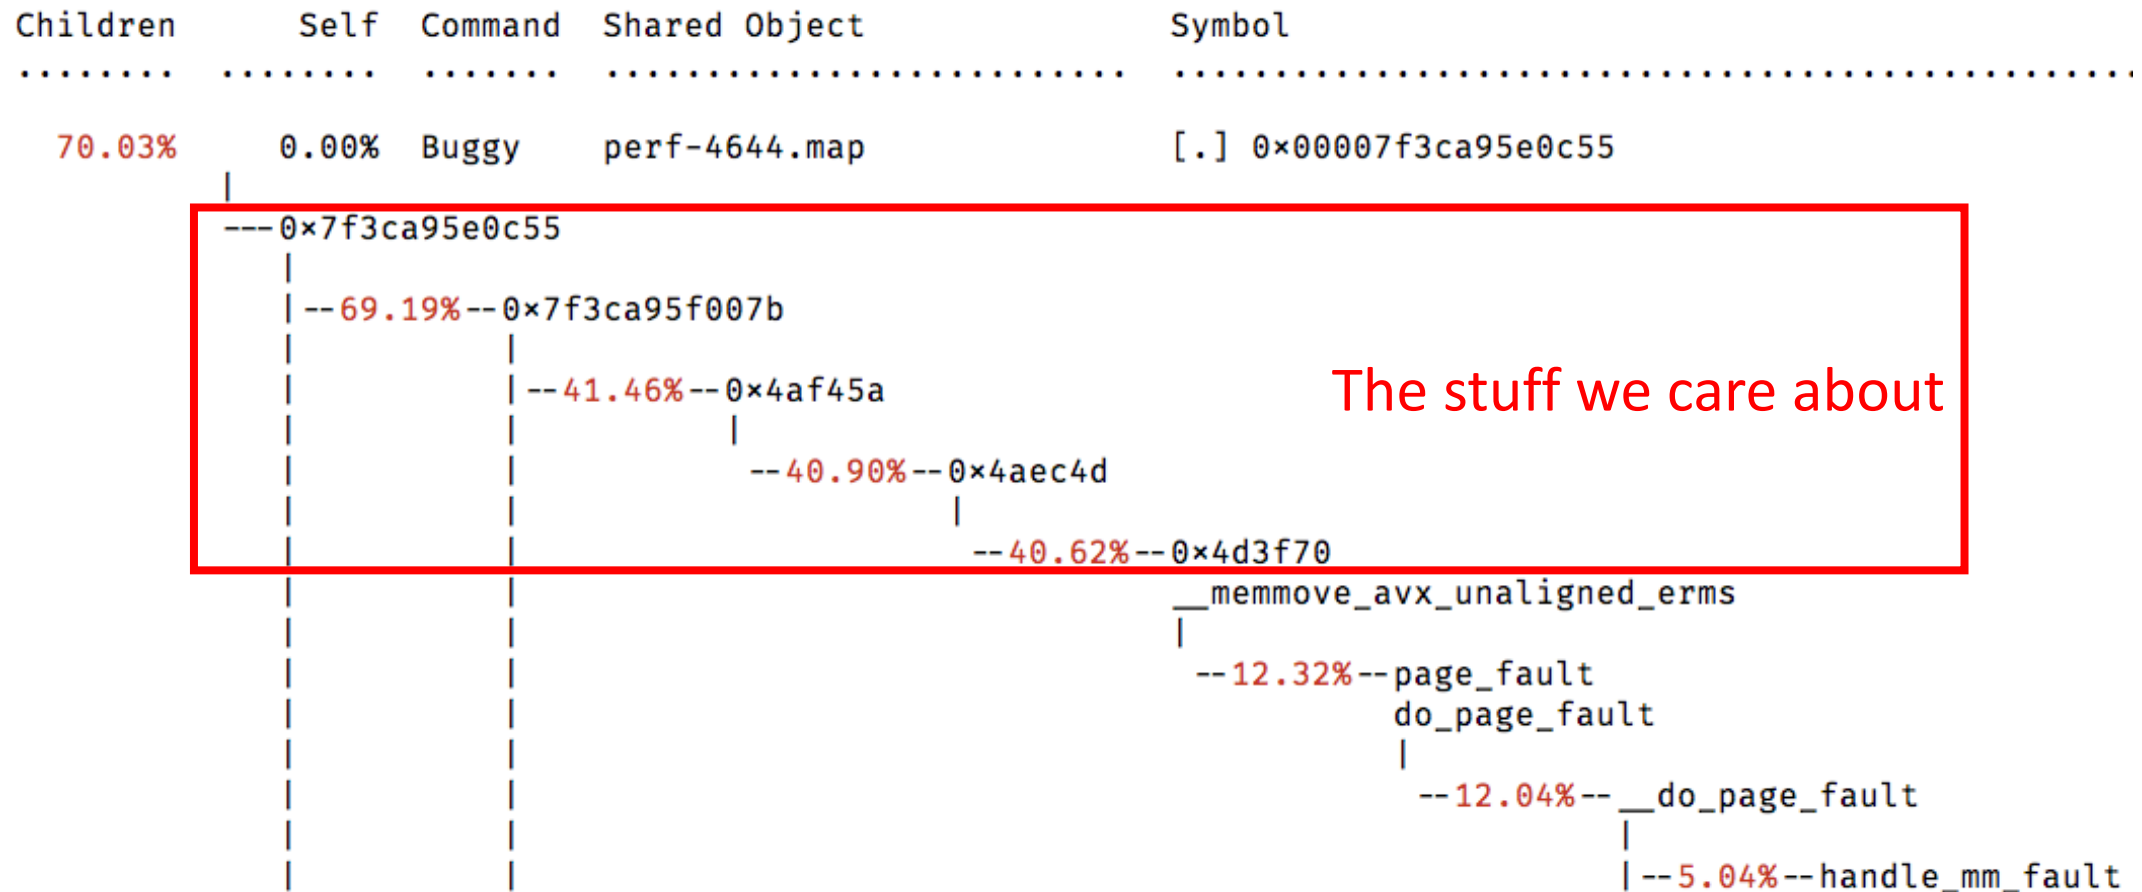
Debian-based: `apt install linux-tools-common`

RedHat-based: `yum install perf`

perf_events Architecture



Five Things That Will Happen To You If You Don't Have Symbolic Debug Information



Getting Debug Information

	Type	Debug information source
SyS_write	Kernel	/proc/kallsyms
__write	Native	Debuginfo package
System.IO.SyncText...	Managed (AOT)	Crossgen*
System.Console.WriteLine	Managed (AOT)	Crossgen*
MyApp.Program.Foo	Managed (JIT)	/tmp/perf-\$PID.map
MyApp.Program.Main	Managed (JIT)	/tmp/perf-\$PID.map
ExecuteAssembly	Native (CLR)	Debuginfo package or source build
CorExeMain	Native (CLR)	Debuginfo package or source build
__libc_start_main	Native	Debuginfo package

Flame Graphs

A visualization method (adjacency graph), very useful for stack traces, invented by Brendan Gregg

<http://www.brendangregg.com/flamegraphs.html>

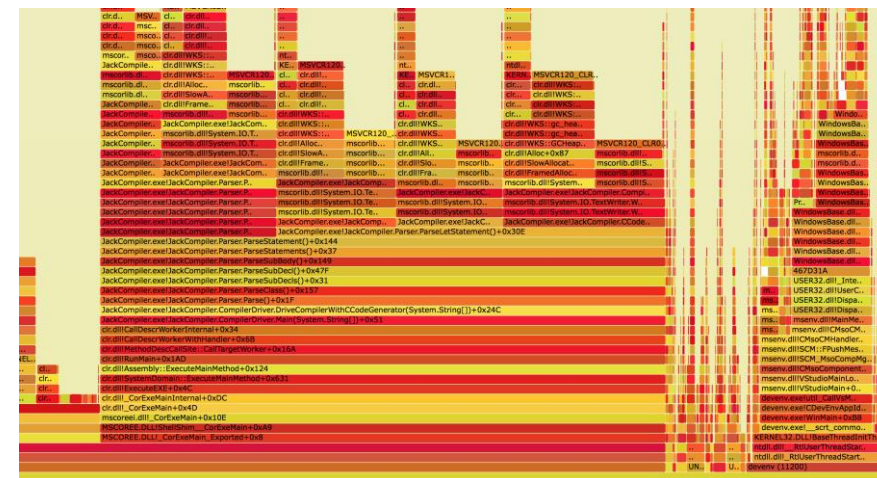
Turns thousands of stack trace pages into a single interactive graph

Example scenarios:

- Identify CPU hotspots on the system/application

- Show stacks that perform heavy disk accesses

- Find threads that block for a long time and the stack where they do it



Reading a Flame Graph

Each rectangle is a function

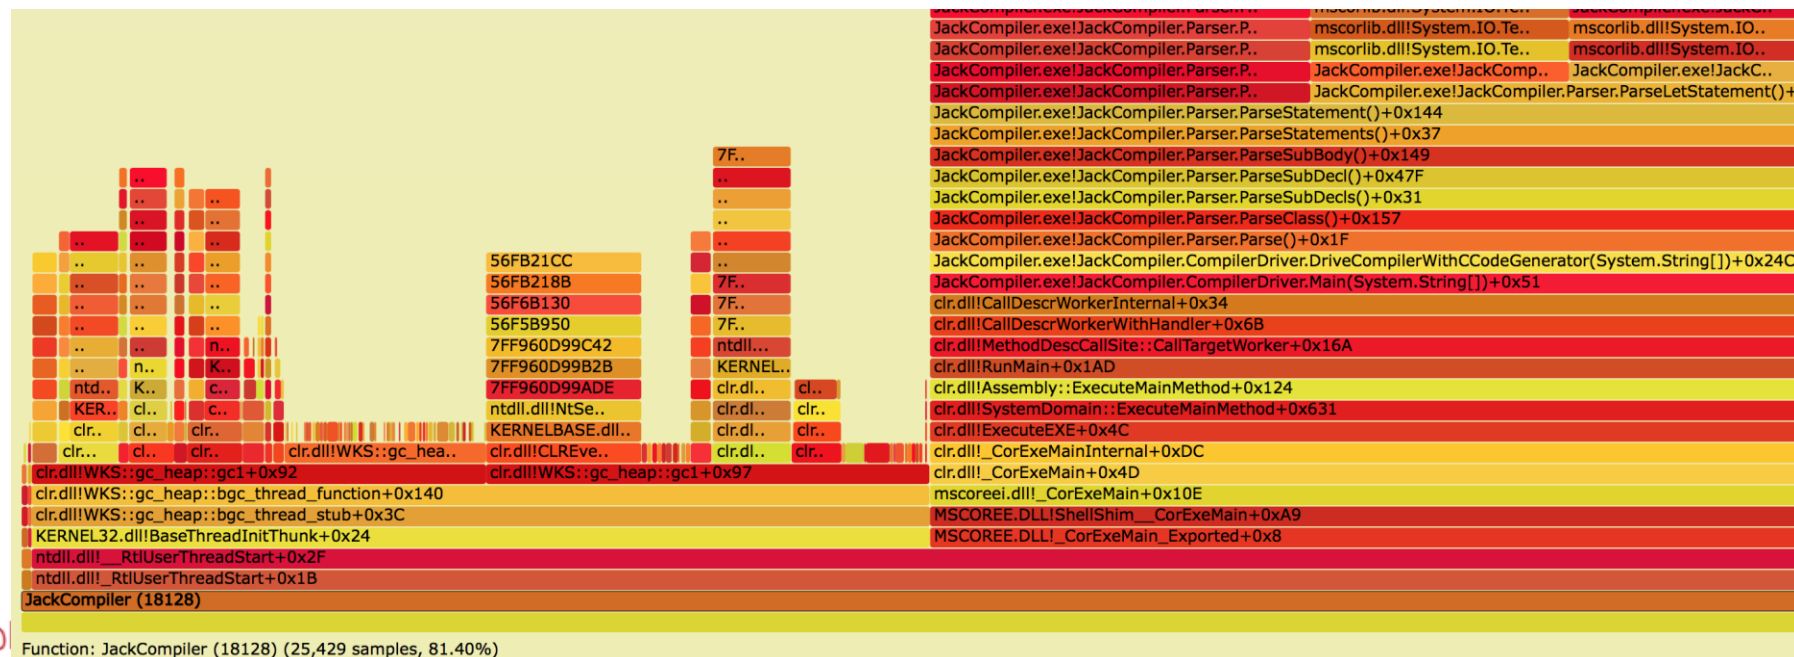
Y-axis: caller-callee

X-axis: sorted stacks (not time)

Wider frames are more common

Supports zoom, find

Filter with grep 🕶️



Demo:

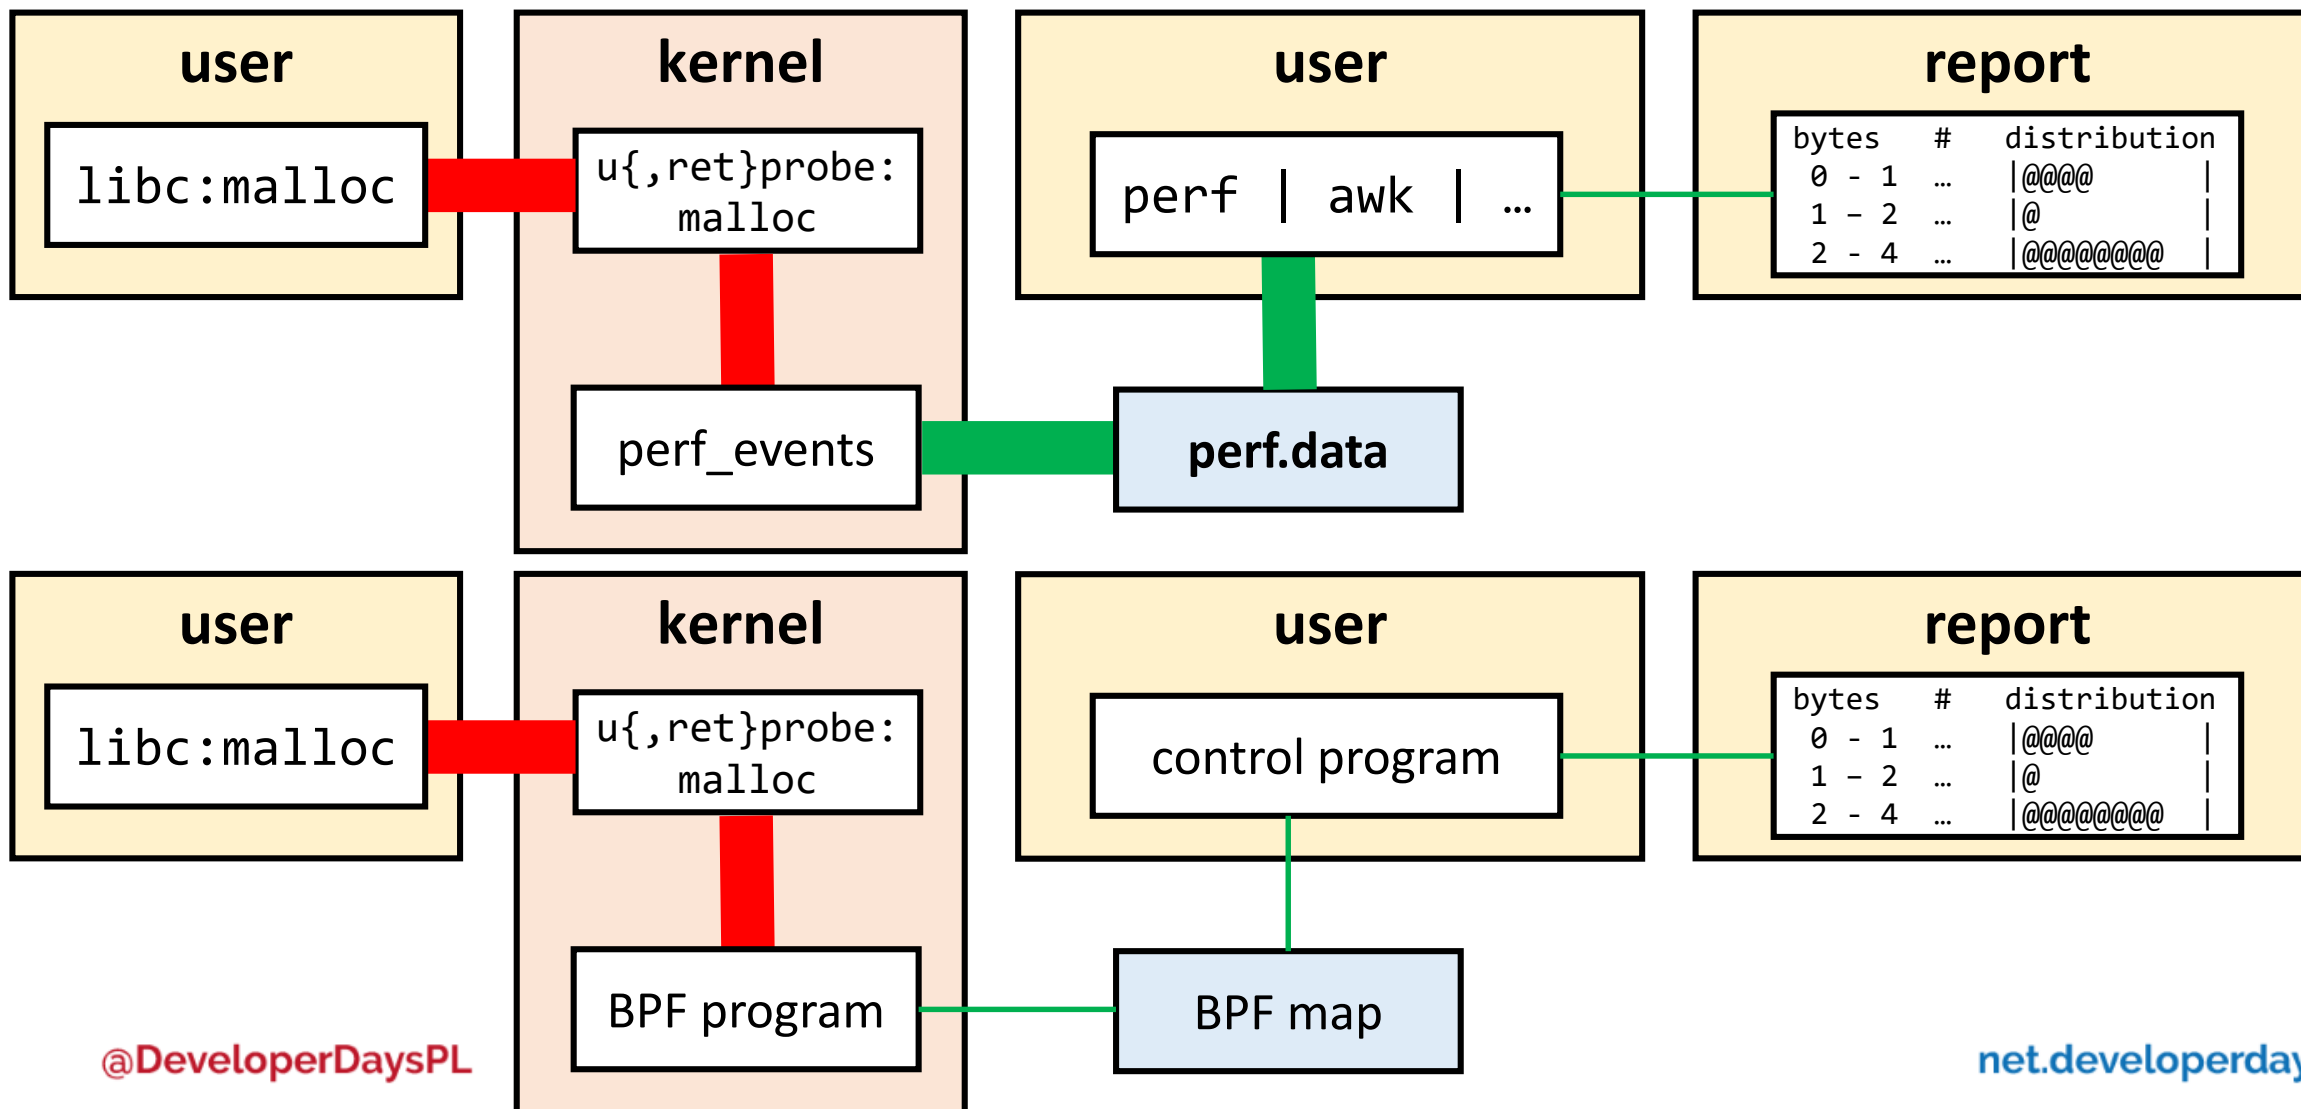
CPU Profiling With Flame Graphs

1\$ make build && make run

2\$ make authbench

3\$ make authrecord

The Old Way And The New Way: BPF



The BCC BPF Front-End

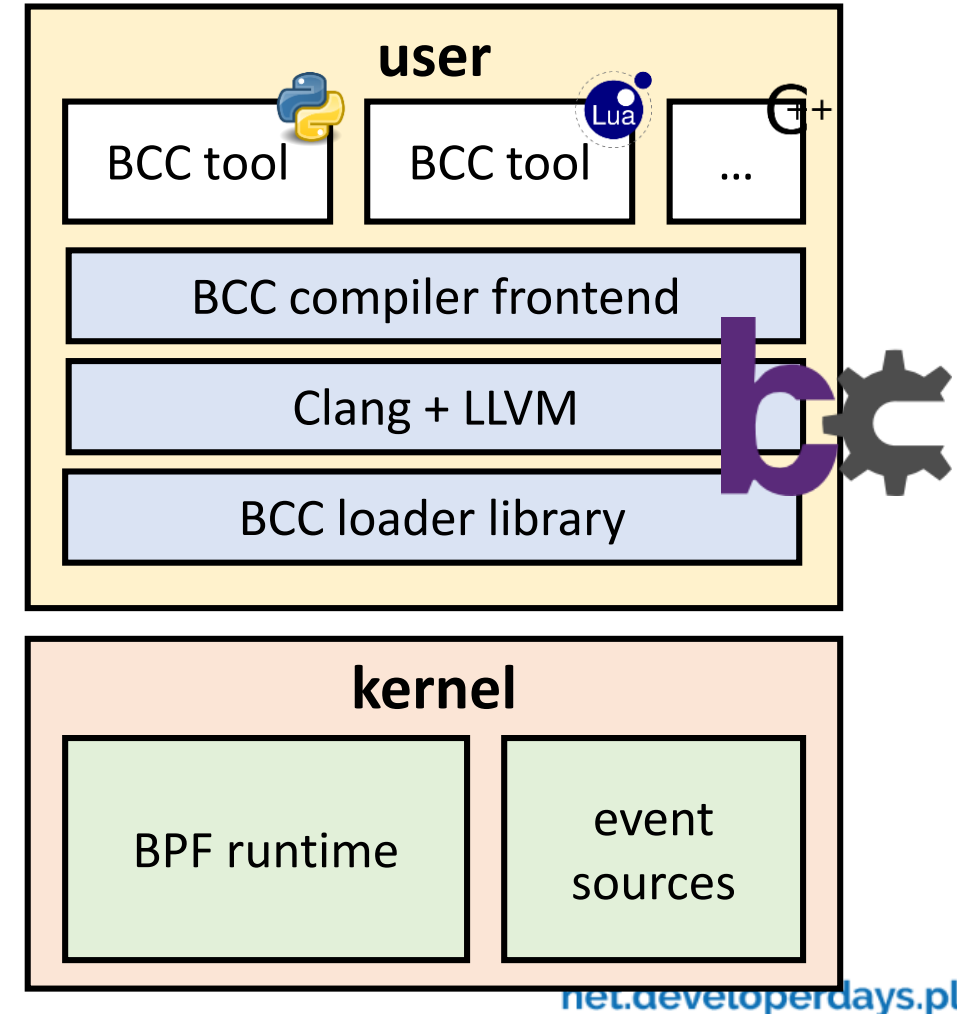
<https://github.com/iovisor/bcc>



BPF Compiler Collection (BCC) is a BPF frontend library and a massive collection of performance tools

Contributors from Facebook, PLUMgrid, Netflix, Sela

Helps build BPF-based tools in high-level languages

Python, Lua, C++



runqlat
cpudist
offcputime
offwaketime
cpuunclaimedmemleak
oomkill
slabratetophardirqs
softirqs
ttsnooptcptop
tcplife
tcpconnect
tcpacceptbiotop
biolateness
biosnoop
bitesizefiletop
filelife
fileslower
vfscount
vfsstat
cachestat
cachetop
mountsnoop
*fsslower
*fsdist
dcstat
dcsnoop
mdflushargdist
trace
funccount
funclatency
stackcountexecsnoop
opensnoop
killsnoop
statsnoop
syncsnoop
setuidsnoopmemleak
sslsniff
gethostlatency
deadlock_detectormysqld_qslower
bashreadline
dbslower
dbstat
mysqldsniffustat
ugc
uthreads
ucalls
uflow

Applications

System libraries

Managed runtimes

Syscall interface

Filesystem

TCP/IP

Scheduler

Mem

Block I/O

Ethernet

Device drivers

Demo:

Tracing .NET Core Apps

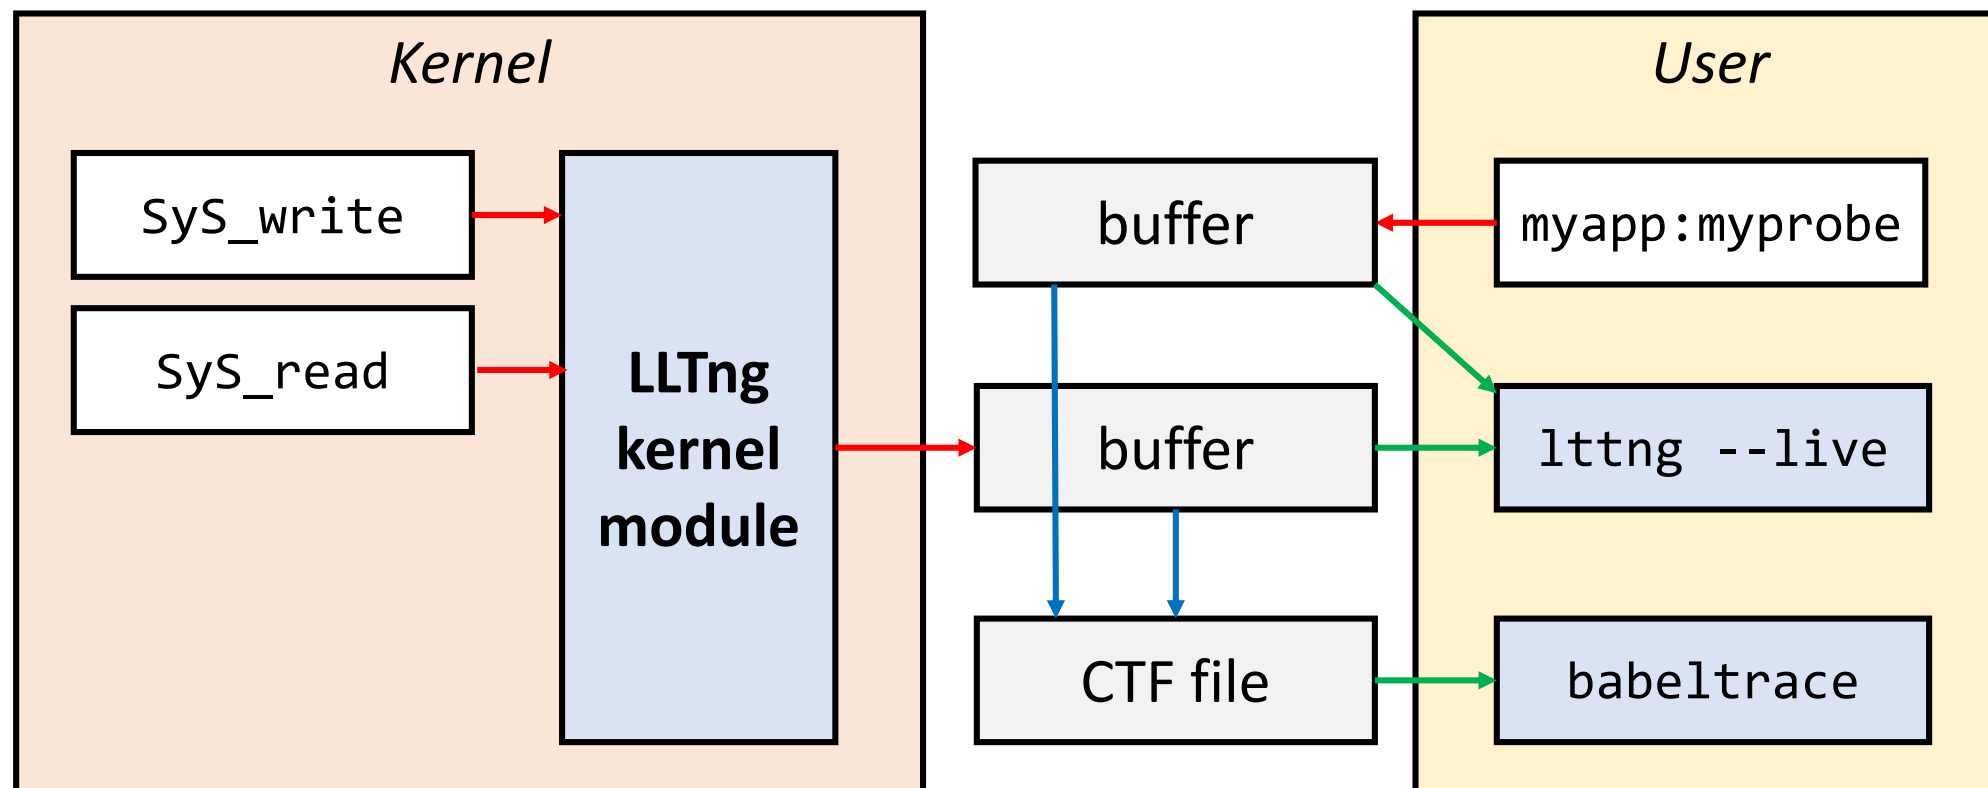
1\$ make build && make run

2\$ make getstatsbench

3\$ make getstatsrecord 3\$ make alloccount 3\$ make allocstacks

2\$ make catsbench 3\$ make catsrecord2

LTtng Architecture



Demo:


Capturing Runtime Events

```
1$ make build && make run
```

```
2$ make getstatsbench
```

```
3$ make gcrecord 3$ make gcview | grep... 3$ make gcallocstats
```


*Plotting data from GCStats events,
[original work](#) by Aleksei Vereshchagin
at DotNext Saint Petersburg*



Core Dumps

A core dump is a memory snapshot of a running process

Can be generated **on crash** or **on demand**



Generating Core Dumps

/proc/sys/kernel/core_pattern configures the core file name or application to process the crash

ulimit -c controls maximum core file size (often 0 by default)

gcore (part of gdb) can create a core dump on demand

Analyzing .NET Core Dumps

```
$ llldb /usr/bin/dotnet -c core.1788
```

```
(lldb) bt
```

```
thread #1: tid = 0, 0x00007f7c3a37c7ef libc.so.6`gsignal + 159, name = 'Buggy', stop reason = signal SIGABRT
* frame #0: 0x00007f7c3a37c7ef libc.so.6`gsignal + 159
  frame #1: 0x00007f7c3a37e3ea libc.so.6`abort + 362
  frame #2: 0x00007f7c399a90bc libcoreclr.so`PROCAbort + 124
  frame #3: 0x00007f7c399a7fbb libcoreclr.so`PROCEndProcess(void*, unsigned int, int) + 235
  frame #4: 0x00007f7c39711318 libcoreclr.so`UnwindManagedExceptionPass1(PAL_SEHException&, _CONTEXT*) + 840
  frame #5: 0x00007f7c397113c9 libcoreclr.so`DispatchManagedException(PAL_SEHException&, bool) + 73
  frame #6: 0x00007f7c39681afa libcoreclr.so`IL Throw(Object*) + 794
  frame #7: 0x00007f7bc05609e2
  frame #8: 0x00007f7bbfff7d0e
  frame #9: 0x00007f7bbfff9349
  frame #10: 0x00007f7c3971da46 libcoreclr.so`FastCallFinalizeWorker + 6
  frame #11: 0x00007f7c395c0c28 libcoreclr.so`MethodTable::CallFinalizer(Object*) + 600
  frame #12: 0x00007f7c396627de libcoreclr.so`FinalizerThread::DoOneFinalization(Object*, Thread*, int, bool*) + 334
  frame #13: 0x00007f7c396625ea libcoreclr.so`FinalizerThread::FinalizeAllObjects(Object*, int) + 266
  frame #14: 0x00007f7c39662c4e libcoreclr.so`FinalizerThread::FinalizerThreadWorker(void*) + 446
  frame #15: 0x00007f7c395fea62 libcoreclr.so`ManagedThreadBase_DispatchOuter(ManagedThreadCallState*) + 402
  frame #16: 0x00007f7c395ff2be libcoreclr.so`ManagedThreadBase::FinalizerBase(void (*)(void*)) + 94
  frame #17: 0x00007f7c39662ecc libcoreclr.so`FinalizerThread::FinalizerThreadStart(void*) + 204
  frame #18: 0x00007f7c399aad22 libcoreclr.so`CorUnix::CPalThread::ThreadEntry(void*) + 306
  frame #19: 0x00007f7c3afb6dca libpthread.so.0`start_thread + 202
  frame #20: 0x00007f7c3a44f0af libc.so.6`clone + 95
```

The stuff we care about

libsosplugin.so

```
(lldb) plugin load ../libsosplugin.so
```

```
(lldb) setclrpath ...
```

DumpObj

DumpArray

DumpHeap

GCRoot

PrintException

Threads

ClrStack

EEHeap

DumpDomain

DumpAssembly

Demo:

Dump Generation And Analysis

1\$ make dockersvc && make dockerrun

2\$ make update

3\$ make updatelogs 3\$ make updateanalyze

Checklist: Preparing Your Environment

- ☐ `export COMPlus_PerfMapEnabled=1`
- ☐ AOT perf map with crossgen
- ☐ Debuginfo package for libcoreclr, libc
- ☐ Install perf/BCC tools
- ☐ `export COMPlus_EnableEventLog=1`
- ☐ `ulimit -c unlimited` (or managed by system)
- ☐ Install gdb (for gcore), lldb-3.x

Summary

We have learned:

- ✓ To profile CPU activity in .NET Core apps
- ✓ To visualize stack traces (e.g. of CPU samples) using flame graphs
- ✓ To use Linux tracing tools with .NET Core processes
- ✓ To capture .NET Core runtime events using LTTng
- ✓ To generate and analyze core dumps of .NET Core apps

References

perf and flame graphs

https://perf.wiki.kernel.org/index.php/Main_Page

<http://www.brendangregg.com/perf.html>

<https://github.com/brendangregg/perf-tools>

.NET Core diagnostics docs

<https://github.com/dotnet/coreclr/blob/master/Documentation/project-docs/linux-performance-tracing.md>

<https://github.com/dotnet/coreclr/blob/master/Documentation/building/debugging-instructions.md>

My blog posts

<http://blogs.microsoft.co.il/sasha/2017/02/26/analyzing-a-net-core-core-dump-on-linux/>

<http://blogs.microsoft.co.il/sasha/2017/02/27/profiling-a-net-core-application-on-linux/>

<http://blogs.microsoft.co.il/sasha/2017/03/30/tracing-runtime-events-in-net-core-on-linux/>

BCC tutorials

<https://github.com/iovisor/bcc/blob/master/docs/tutorial.md>

https://github.com/iovisor/bcc/blob/master/docs/tutorial_bcc_python_developer.md

https://github.com/iovisor/bcc/blob/master/docs/reference_guide.md

Thank You!



Slides: <https://s.sashag.net/ddwaw17-3>

Demos & labs: <https://github.com/goldshtn/linux-tracing-workshop>

Sasha Goldshtein
CTO, Sela Group



Please rate this session using



**Developer Days
Mobile App**

or at



login.developerdays.pl

Strategic Sponsors



Gold Sponsors



Silver Sponsors

