

Leveraging KVM as a Debugging Platform



Wenzel



mtarral



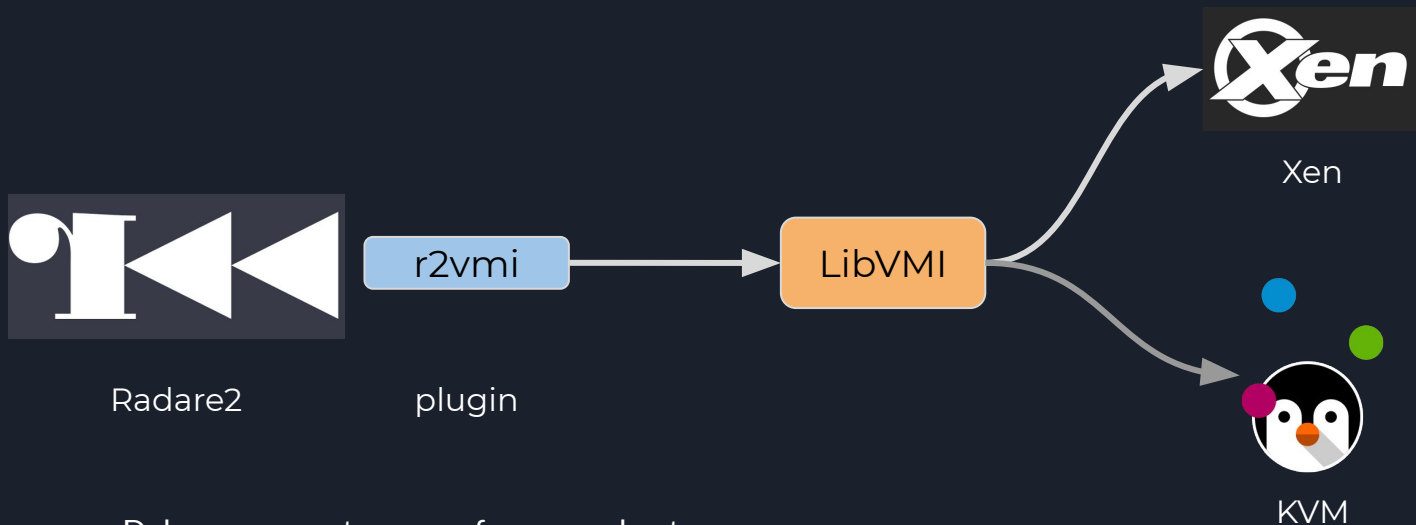
mathieu.tarral



Agenda

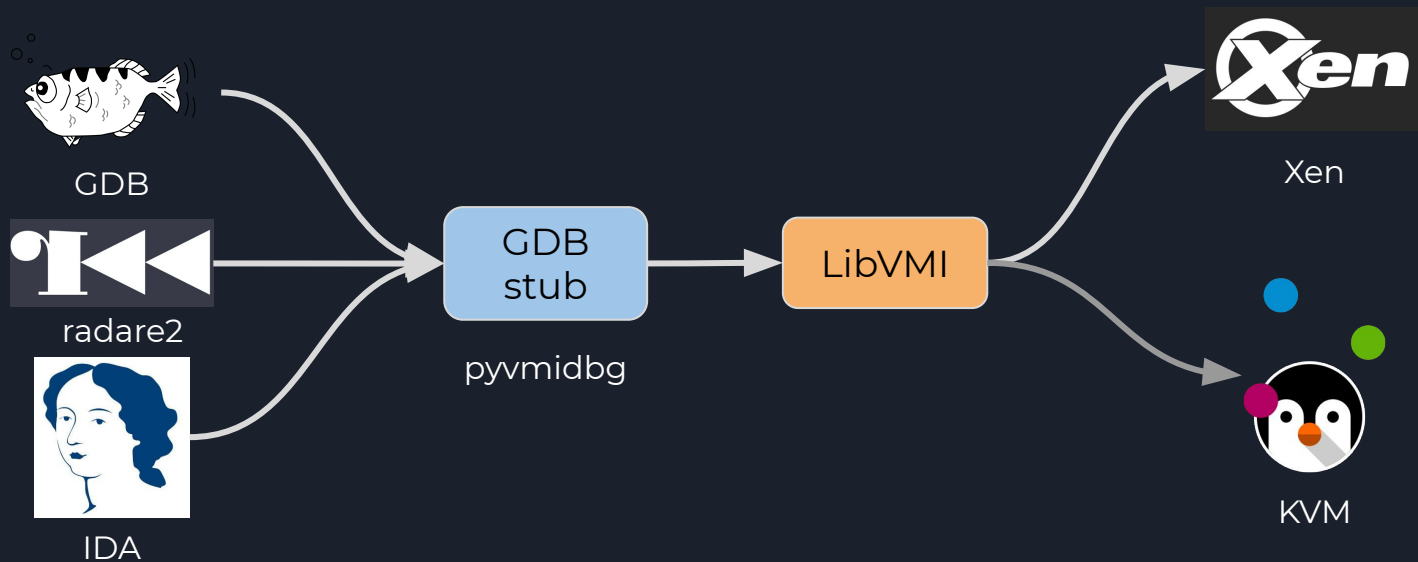
1. VM introspection on KVM
2. KVM-VMI Setup
3. Integration in LibVMI (demo)
4. KVM as a debugging platform (demo)
5. Future

Hack.lu 2018: r2vmi on Xen



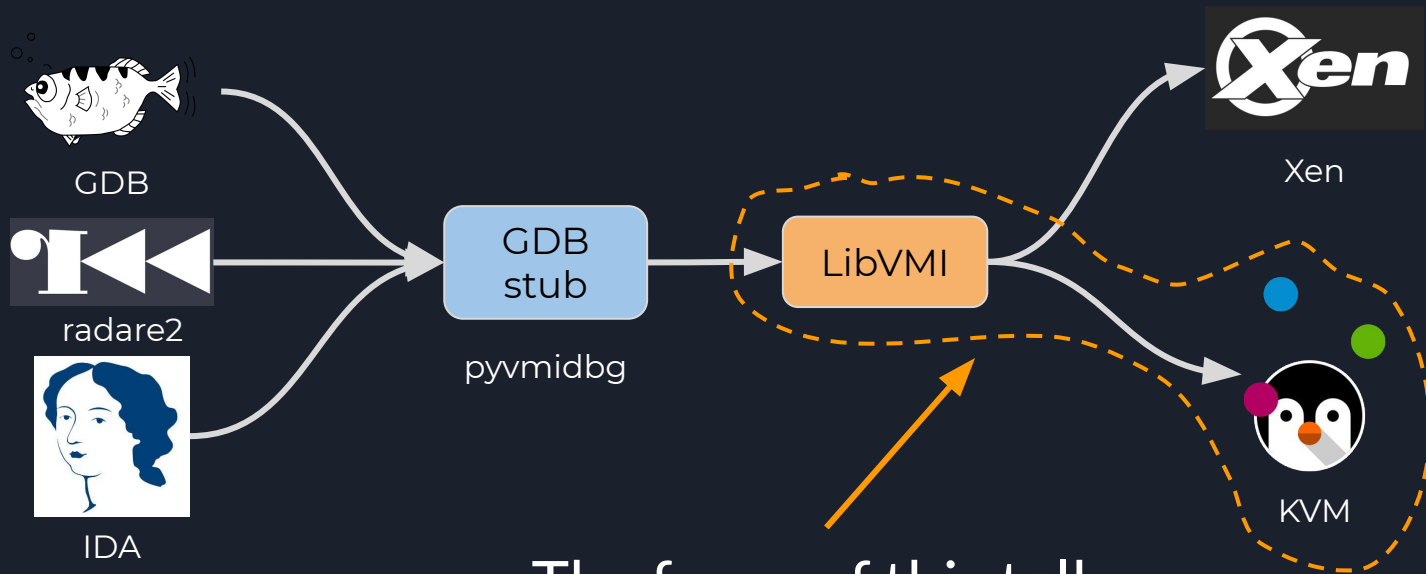
- Debug any guest process from your host
- Radare2 as main tool
- Hypervisor-agnostic by design
 - KVM was not available (missing APIs)

2019: a Python VMI-GDB stub



<https://github.com/Wenzel/pyvmidbg>

Hack.lu 2019: a VMI-GDB stub... on KVM ?!




The focus of this talk:
a new LibVMI **KVM driver**


VM Introspection on KVM ?



VMI API: Hypervisor Support



Patches available
(community)




Upstream integration




Alternate EPT/RVI
available


VMI API: Hypervisor Support



VMI API: Hypervisor Support



Patches available
(community)

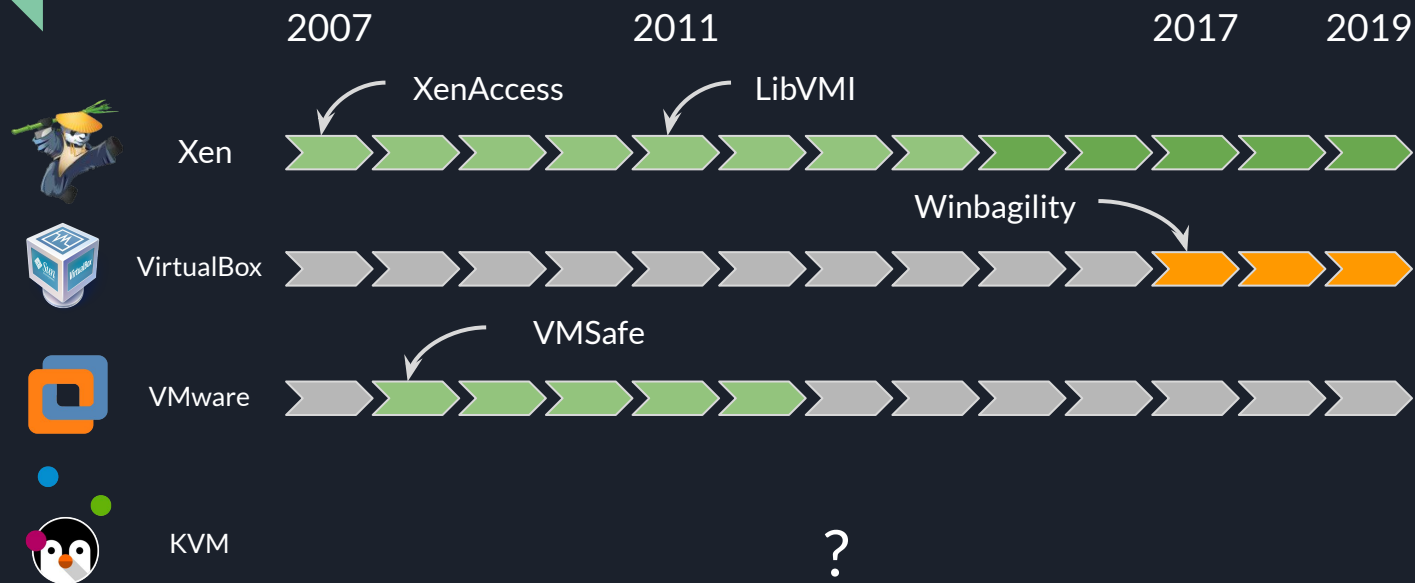


Upstream integration



Alternate EPT/RVI
available

VMI API: Hypervisor Support

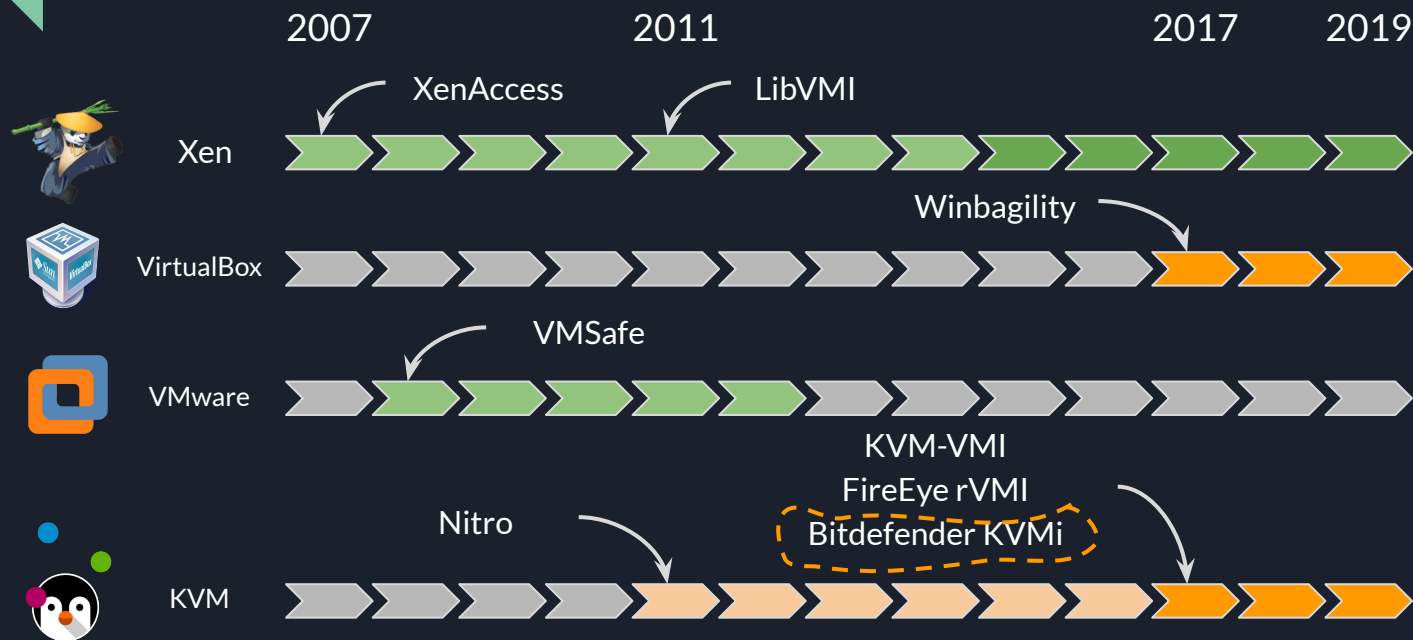


VMI on KVM ?



- VMI-based frameworks are available for QEMU (**full-emulation**)
 - PANDA - 2011
 - DECAF - 2014
 - PyREBox - 2017
- What about **KVM** ?
 - Nitro - 2011
 - syscall interception
 - KVM-VMI community (Github) - Jan 2017
 - improved Nitro with Python framework
 - FireEye rVMI - July 2017
 - full system debugger
 - Bitdefender **KVMi** subsystem - June 2017
 - complete API for VMI
 - KVM-VMI **integrates** BitDefender KVMi patches - January 2018
 - + new LibVMI KVM driver !

VMI API: Hypervisor Support



Patches available
(community)



Upstream integration



Alternate EPT/RVI
available

VMI on KVM: BitDefender KVMi Evolution

v1
June 2017
4.12-rc5

v4
December 2018
4.15-rc2

v5
December 2018
4.20-rc7

v6
August 2019
5.0.0-rc7

- initial API design
- RFC on KVM mailing list
- QEMU will connect to the introspection tool, and perform a handshake
- change mapping protocol
- improved remote-mapping
- single-step support as reply for #PageFault events
- new ioctl to remove hooks when a domain is suspended/live-migrated
- lots of fixes
- speed improvements
- guests are much more stable

Next: drop RFC and **merge** basic introspection in KVM !

Note: v2 and v3 were only documentation updates



KVMi API: Overview

- Get VM hardware state
 - r/w physical memory
 - r/w VCPU registers
 - get domain info: { VPCU count, Max PFN, ... }
 - pause/resume domain
- Listen for hardware events
 - CR0/CR3/CR4
 - MSR
 - interrupts (int3)
 - memory access
 - descriptor
 - hypercall
- Utilities
 - guest remote memory mapping
 - exception injection
 - page fault injection



<https://github.com/KVM-VMi/kvm/blob/kvmi/tools/kvm/kvmi/include/kvmi/libkvmi.h>



Why: BitDefender KVMi

*“At the moment, the target audience for KVMi are **security software authors** that wish to perform **forensics** on newly discovered threats (exploits) or to implement **another layer of security** like preventing a large set of kernel rootkits simply by **"locking" the kernel** image in the shadow page tables (ie. **enforce** .text r-x, .rodata rw- etc.). ”*

KVM-VMi Setup



KVM-VMi Setup - 1

Setup Wiki page: <https://github.com/KVM-VMi/kvm-vmi/wiki/KVM-VMi-setup>

```
git clone https://github.com/KVM-VMi/kvm-vmi.git --recursive --branch kvmi
```

KVM

```
cd kvm
make menuconfig # (CONFIG_KVM_INTROSPECTION)
make bzImage && make modules
sudo make modules_install
sudo make install
sudo reboot
uname -r # (5.0.0-rc7)
```

QEMU

```
cd qemu
./configure --target-list=x86_64-softmmu
make
sudo make install # /usr/local/bin/qemu-system-x86_64
/usr/local/bin/qemu-system-x86_64 --version # 2.11.93
```

KVM-VMi Setup - 2

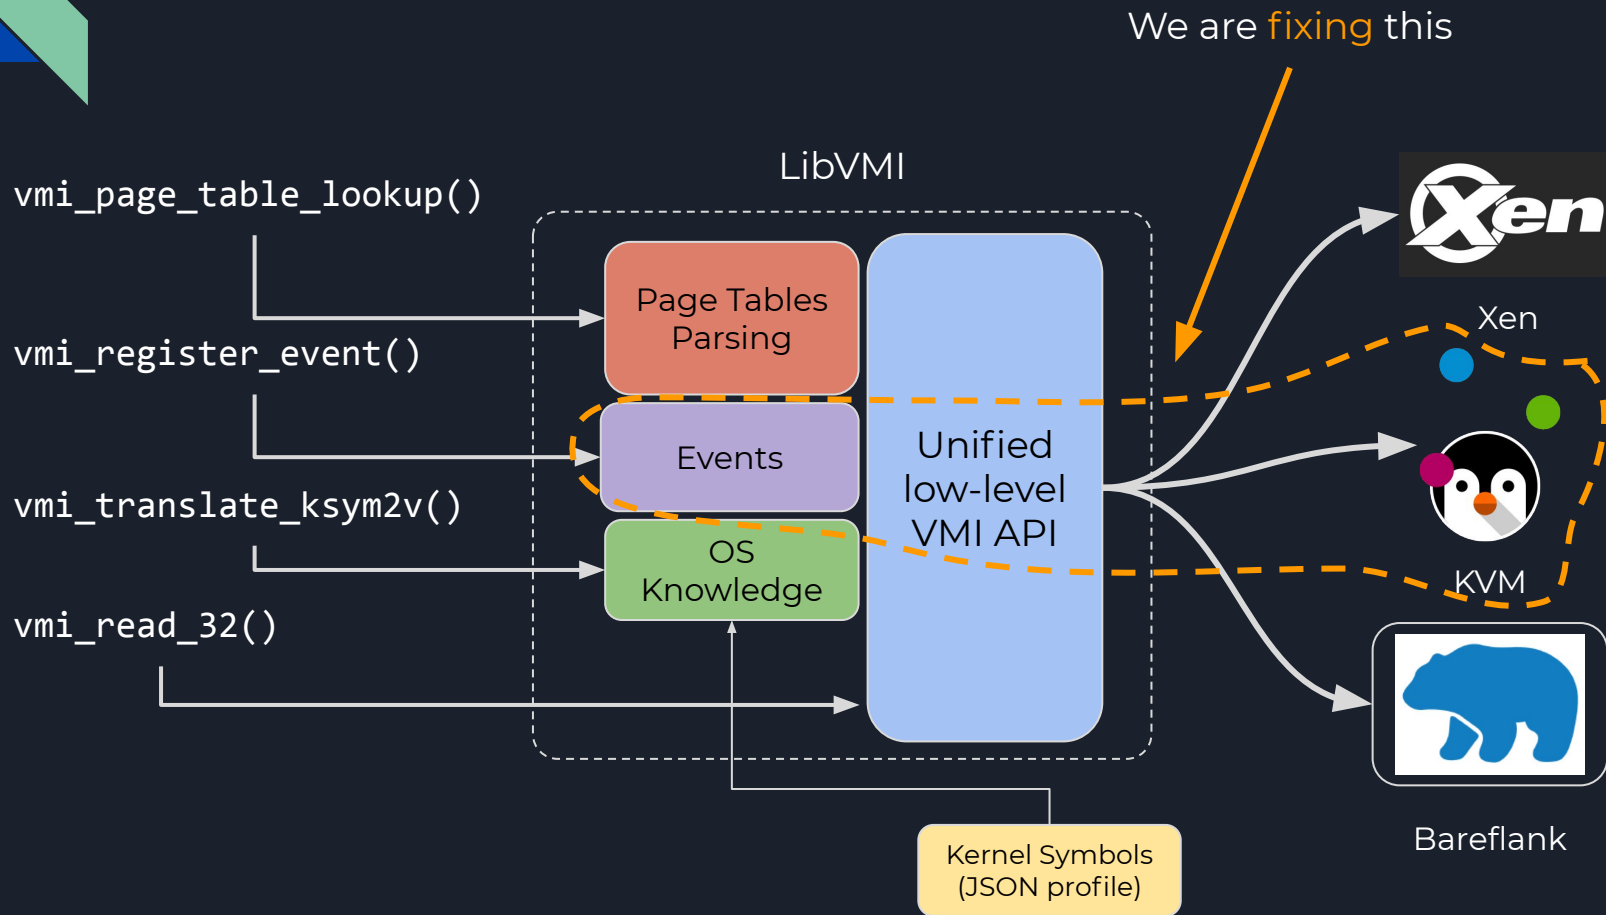
- QEMU has new command line parameters: a `socket` (ex: `/tmp/kvmi_win7.sock`)

```
<domain type='kvm' xmlns:qemu='http://libvirt.org/schemas/domain/qemu/1.0'>
  <qemu:commandline>
    <qemu:arg value='-chardev' />
    <qemu:arg value='socket,path=/tmp/kvmi_win7.sock,id=chardev0,reconnect=10' />
    <qemu:arg value='-object' />
    <qemu:arg value='secret,id=key0,data=some' />
    <qemu:arg value='-object' />
    <qemu:arg value='introspection,id=kvmi,chardev=chardev0,key=key0' />
    <qemu:arg value='-accel' />
    <qemu:arg value='kvm,introspection=kvmi' />
  </qemu:commandline>
  ...
  <devices>
    <emulator>/usr/local/bin/qemu-system-x86_64</emulator>
```

Integration in LibVMI



LibVMI: Overview



LibVMI: Python Bindings - CR3 Events

```
1 # 1 - init LibVMI
2 kvm_socket = {VMIInitData.KVMI_SOCKET: "/tmp/introspector"}
3 with Libvmi("winxp", INIT_DOMAINNAME | INIT_EVENTS, init_data=kvm_socket,
4     partial=True) as vmi:
5     counter = Counter()
6
7     # 2 - define CR3 callback
8     def cr3_callback(vmi, event):
9         cr3_value = event.value
10        logging.info("CR3 change: %s", hex(cr3_value))
11        counter[hex(cr3_value)] += 1
12
13    # 3 - define and register CR3-write event
14    with pause(vmi):
15        # register CR3-write event
16        reg_event = RegEvent(X86Reg.CR3, RegAccess.W, cr3_callback)
17        vmi.register_event(reg_event)
18
19    # 4 - listen for events
20    for i in range(0, 100):
21        vmi.listen(500)
22        logging.info(counter)
```

LibVMI: Python Bindings - Memory Events



```
1 with Libvmi(vm_name, INIT_DOMAINNAME | INIT_EVENTS, init_data=kvm_socket,
  partial=True) as vmi:
2     # 1 - init paging to translate virtual addresses
3     vmi.init_paging(0)
4     with pause(vmi):
5         # 2 - get current RIP -> Guest Frame Number
6         # get current RIP on VCPU 0
7         rip = vmi.get_vcpureg(X86Reg.RIP.value, 0)
8         # get DTB
9         cr3 = vmi.get_vcpureg(X86Reg.CR3.value, 0)
10        dtb = cr3 & ~0xfff
11        # get paddr
12        paddr = vmi.pagetable_lookup(dtb, rip)
13        gfn = paddr >> 12
14
15        # 3 - define mem-event callback
16        def cb_mem_event(vmi, event):
17            logging.info("Mem event at RIP: %s, frame: %s, offset: %s, permissions:
18                %s",
19                        hex(event.x86_regs.rip), hex(event.gla), hex(event.offset),
20                        event.out_access.name)
21
22        # 4 - define and register mem-event
23        mem_event = MemEvent(MemAccess.X, cb_mem_event, gfn=gfn)
24        vmi.register_event(mem_event)
25
26        # 5 - listen
27        while not interrupted:
28            vmi.listen(500)
29        logging.info("stop listening")
```

LibVMI: Python Bindings - MSR Events

```
1 with Libvmi(vm_name, INIT_DOMAINNAME | INIT_EVENTS, init_data=kvm_socket,
   partial=True) as vmi:
2
3     # define MSR callback
4     def msr_callback(vmi, event):
5         logging.info("%s %s = %s", name, hex(event.msr), hex(event.value))
6
7     # define and register MSR event
8     with pause(vmi):
9         reg_event = RegEvent(MSR.ALL, RegAccess.W, msr_callback)
10        vmi.register_event(reg_event)
11
12    # listen for events
13    logging.info("listening")
14    while not interrupted:
15        vmi.listen(500)
```

LibVMI: Python Bindings - Demo

Click Me,
I'm a video



LibVMi: Syscall Entrypoint OS Hardening

```
FFDFF20B hook_sysenter proc near ; CODE XREF: seg000:jump_hook_msr↑j
FFDFF20B nop
FFDFF20C pop ebx
FFDFF20D mov ecx, 176h ; msr[0x176]: IA32_SYSENTER_EIP
FFDFF212 rdmsr
FFDFF214 mov ds:orig_msr_value, eax
FFDFF219 lea eax, [ebx+17h] ; FFDFF20A+17=FFDFF221
FFDFF219 ; -- FFDFF221 will become start of the
FFDFF219 ; SYSENTER routine hook
FFDFF21C xor edx, edx
FFDFF21E wrmsr
FFDFF220 retn
```


```
1 msr_counter = Counter()
2 def msr_callback(vmi, event):
3     logging.info("%s %s = %s", name, hex(event.msr), hex(event.value))
4     msr_counter[event.msr] += 1
5     # IA32_SYSENTER_EIP written twice ??
6     if msr_counter[0x176] > 1:
7         # EternalBlue exploit
8         # kill current process
```

KVM as a debugging platform



pyvmidbg on KVM: Windows 10 x64

Click Me,
I'm a video





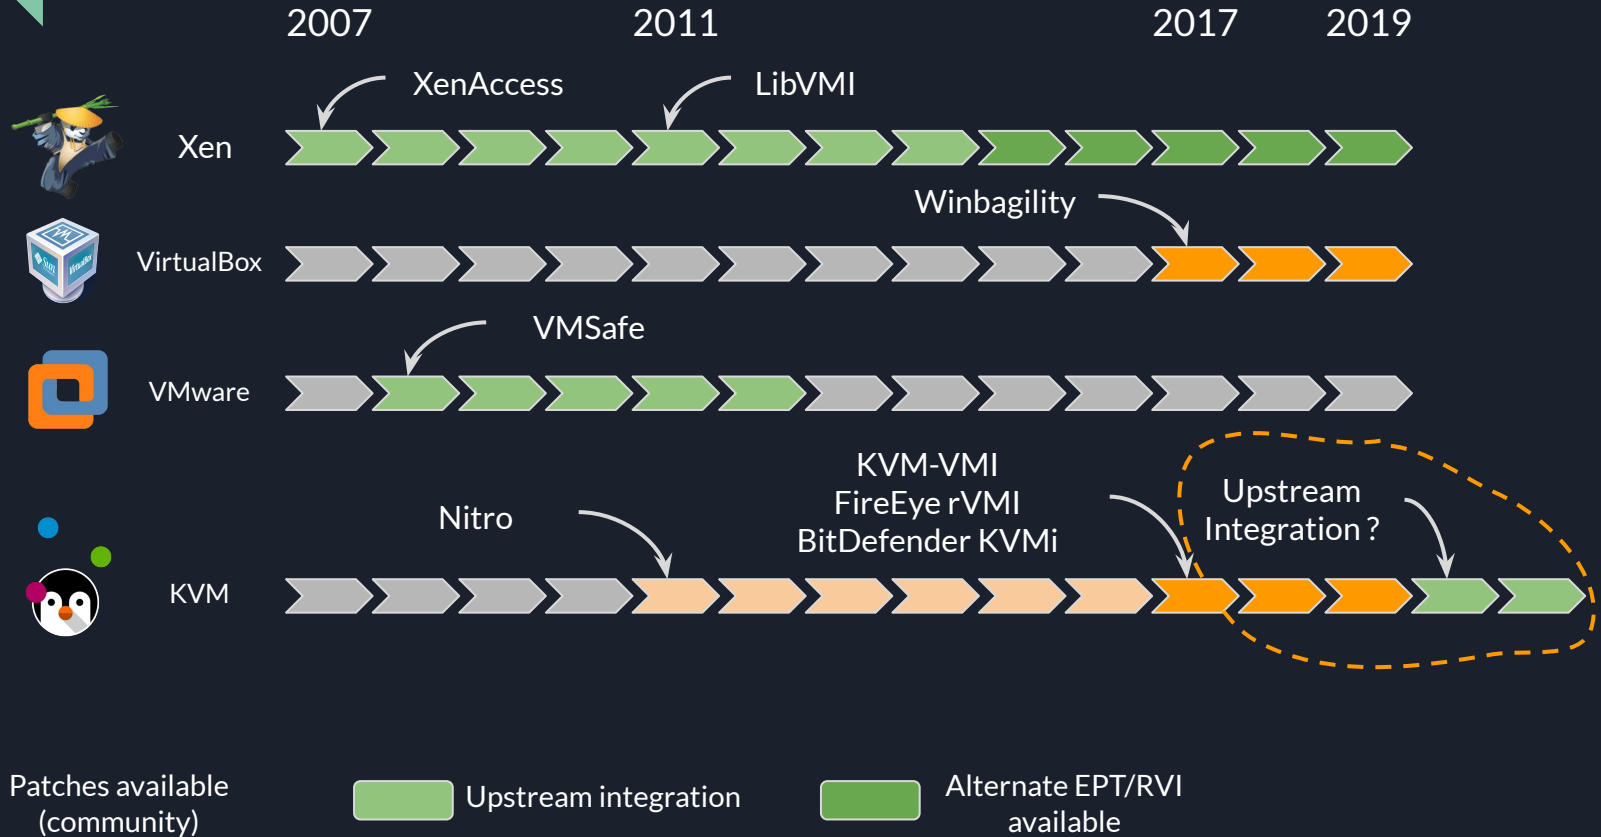
Room for Improvements

- Entrypoint
 - already implemented, but lacks proper pagefault injection to work
 - break on `KiUserThreadStart`
 - check process name
 - read and break on first `ETHREAD.StartAddress` (`ntdll!RtlUserThreadStart`)
 - read and break on entrypoint (`pfnStartAddr` parameter)
- Pagefaults
 - API is available in LibVMI/KVMi, used in pyvmidbg, but not working yet
- Symbols loading
 - enumerate loaded libraries
 - get full image path
 - run libguestfs instance to dynamically copy the binary on the host !
- Stealth breakpoints
 - guard int3 by memory access events, return fake data
 - alternate SLAT API in KVMi ?
- Add more Stubs
 - KD
 - LLDB

Future



VMI API: Hypervisor Support



Introspection on KVM: Use Cases



KVM-based

Debugging

Malware Analysis

Live-Memory Analysis

OS Hardening


Monitoring

Fuzzing


Towards a High-Level VMI Abstraction Library

VMI Apps

Debugging
Malware Analysis
Live-Memory Analysis
OS Hardening
Monitoring
Fuzzing



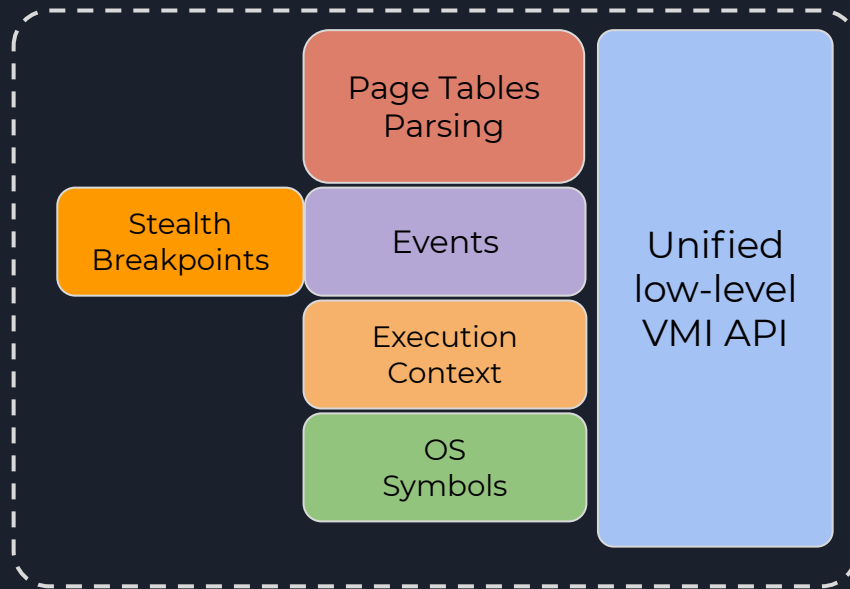
Hypervisors



Towards a High-Level VMI Abstraction Library

VMI Apps

Debugging
Malware Analysis
Live-Memory Analysis
OS Hardening
Monitoring
Fuzzing



Emulators



Custom Hypervisor

Hypervisors

Towards a High-Level VMI Abstraction Library

VMI Apps

Dynamic Analysis

- pyvmidbg
- icebox
- rVMI
- LiveCloudKd
- DECAF
- PANDA
- PyREBox
- Drakvuf

Live-Memory Analysis

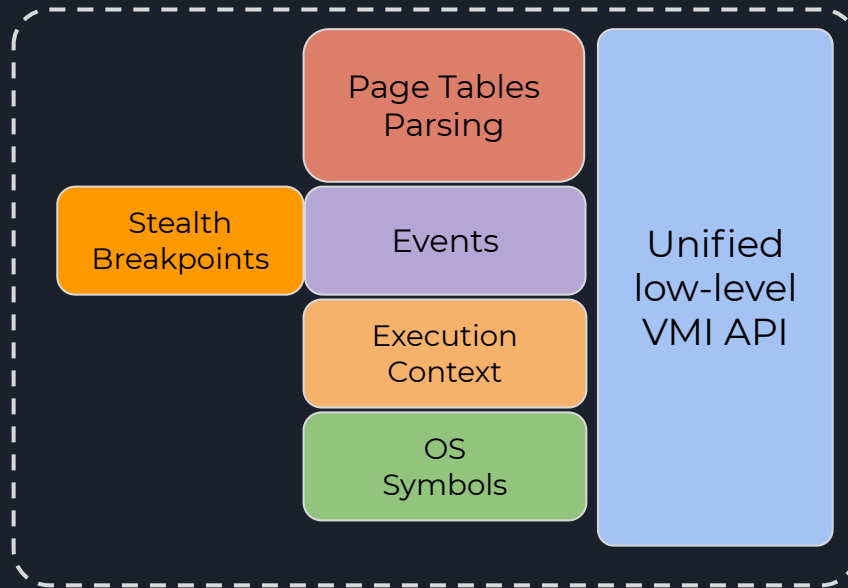
- Volatility
- Rekall

OS Hardening


Monitoring

Fuzzing

- ApplePie



Emulators



Custom Hypervisor

Hypervisors

High-Level VMI Library: What about Rust ?



<https://github.com/Wenzel/libmicrovmi>

VMI Apps

Dynamic Analysis

- pyvmidbg
- icebox
- rVMI
- LiveCloudKd
- DECAF
- PANDA
- PyREBox
- Drakvuf

Live-Memory Analysis

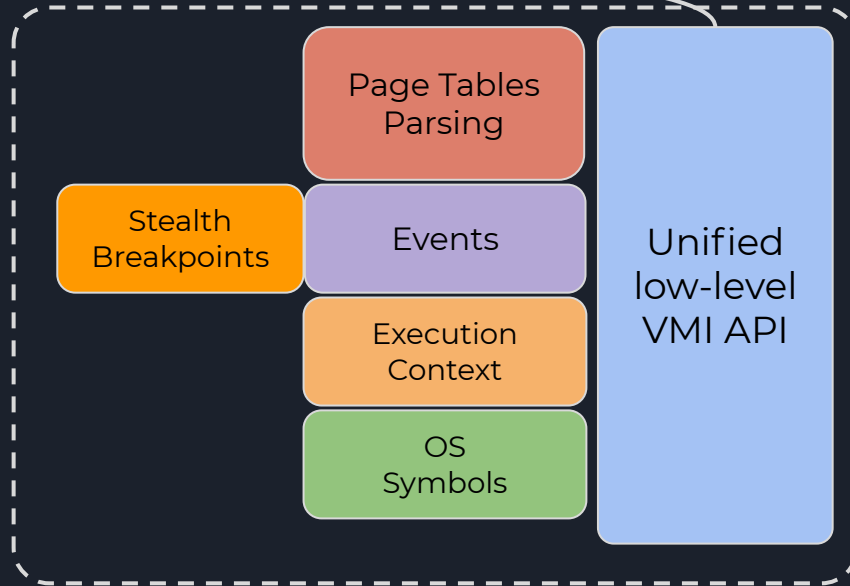
- Volatility
- Rekall

OS Hardening


Monitoring

Fuzzing

- ApplePie



Emulators



Custom Hypervisor

Hypervisors

Libmicrovmi : API example



```
1 // select drive type (Xen, KVM, ...)
2 let drv_type = DriverType::KVM;
3 // init library
4 let mut drv: Box<dyn Introspectable> = microvmi::init(drv_type, domain_name);
5 // pause VM
6 drv.pause().expect("Failed to pause VM");
7 // get max physical address
8 let max_addr = drv.get_max_physical_addr()
9                 .expect("Failed to get max physical address");
10 // read physical memory
11 let mut buffer: [u8; 4096] = [0; 4096];
12 let result = drv.read_physical(0x804d7000, &mut buffer);
13 // resume VM
14 drv.resume().expect("Failed to resume VM");
```



Conclusion

1. Native **introspection** on **KVM** is becoming a reality
2. A new LibVMI **KVM** driver is available to fully exploit its capabilities
3. **KVM**-based full-system debugging is just one of many possibilities
4. Building a **high-level**, **cross-platform**, **multi-hypervisor** VMI abstraction library is a condition to let the VMI ecosystem grow and mature



Thanks

- Mihai Dontu & Adalbert Lazar (Bitdefender)
 - adding singlestep support 35 days before the talk ;)
- Hady Azzam (@hady_azzam)
 - adding Linux support to pyvmidbg !
- Petr Beneš
- Tamas K. Lengyel
- Hack.lu team !

Leveraging KVM as a Debugging Platform



[KVM-VMI/kvm-vmi](https://github.com/KVM-VMI/kvm-vmi)

[Wenzel/pyvmdbg](https://github.com/Wenzel/pyvmdbg)

[Wenzel/libmicrovmi](https://github.com/Wenzel/libmicrovmi)



Wenzel



mtarral



mathieu.tarral



Annex: Experimental Build

- The demos in this presentation have been made on an experimental build
 - KVM-VM/ kvm: <https://github.com/adlazar/kvm/tree/kvmi-v6-single-step>
 - for singlestep events support
 - mtaral/libvmi: https://github.com/mtaral/libvmi/tree/kvmi_events_v6_sstep
 - for singlestep event support in LibVMI KVM
 - KVM-VM/ python:
 - init_data: https://github.com/KVM-VM/python/tree/init_data
 - to init LibVMI Python with a kvmi socket parameter
 - msr: <https://github.com/libvmi/python/pull/48>
 - add support for MSR registers
 - Wenzel/pyvdbg:
 - kvm_support: <https://github.com/Wenzel/pyvdbg/pull/40>
 - add KVM support